



Chapter 3

Processes

Da-Wei Chang

CSIE.NCKU

Source: Abraham Silberschatz, Peter B. Galvin, and Greg Gagne, "Operating System Concepts", 10th Edition, Wiley.

Outline




- Process Concept
- Process Scheduling
- Operations on Processes
- Inter-Process Communication (IPC)
- Examples of IPC Systems
- Communication in Client-Server Systems

Process Concept

- An operating system executes user programs
 - Batch system – jobs
 - Time-sharing systems – processes or tasks
- We use the terms *job* and *process* almost interchangeably
- Process – a program in execution; process execution must progress in sequential fashion
- A process includes
 - text (i.e., code section)
 - data section
 - heap
 - stack
 - program counter and the content of the processor registers
- Program – *passive* ; Process – *active*

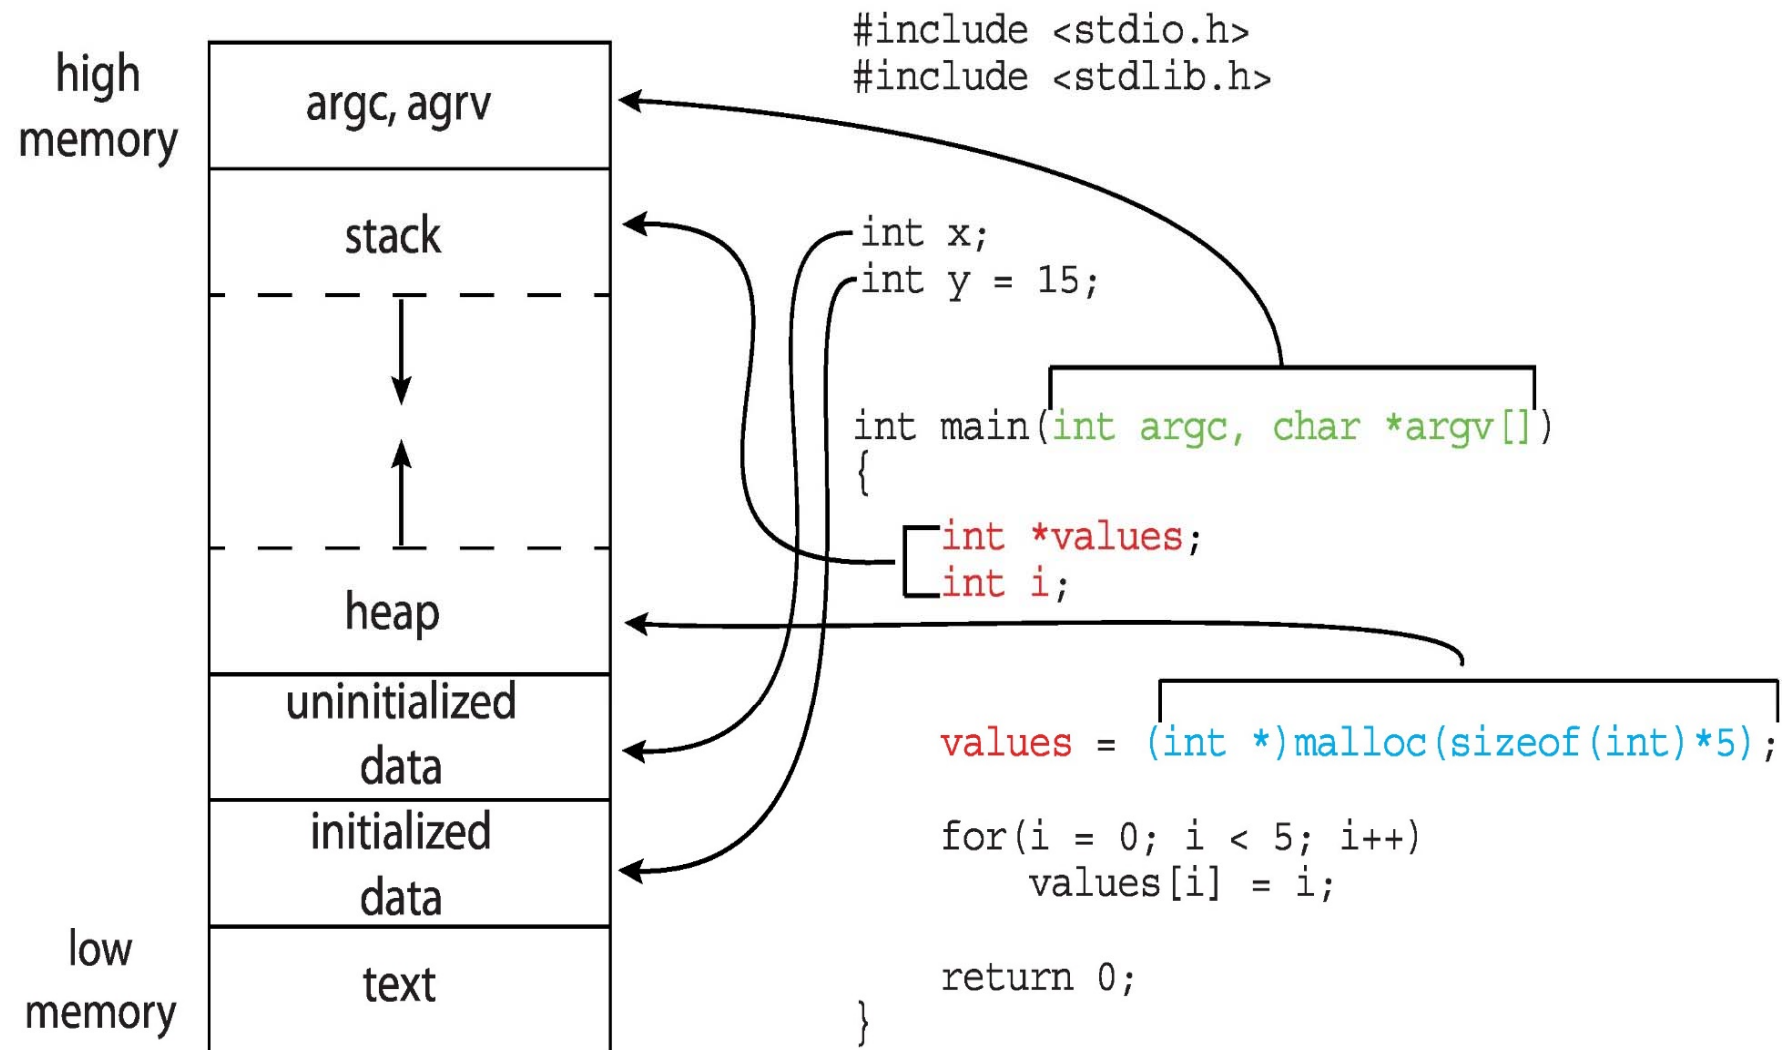
A Process in Memory



More than one processes can be associated with the same program

Which parts are the same for these processes?

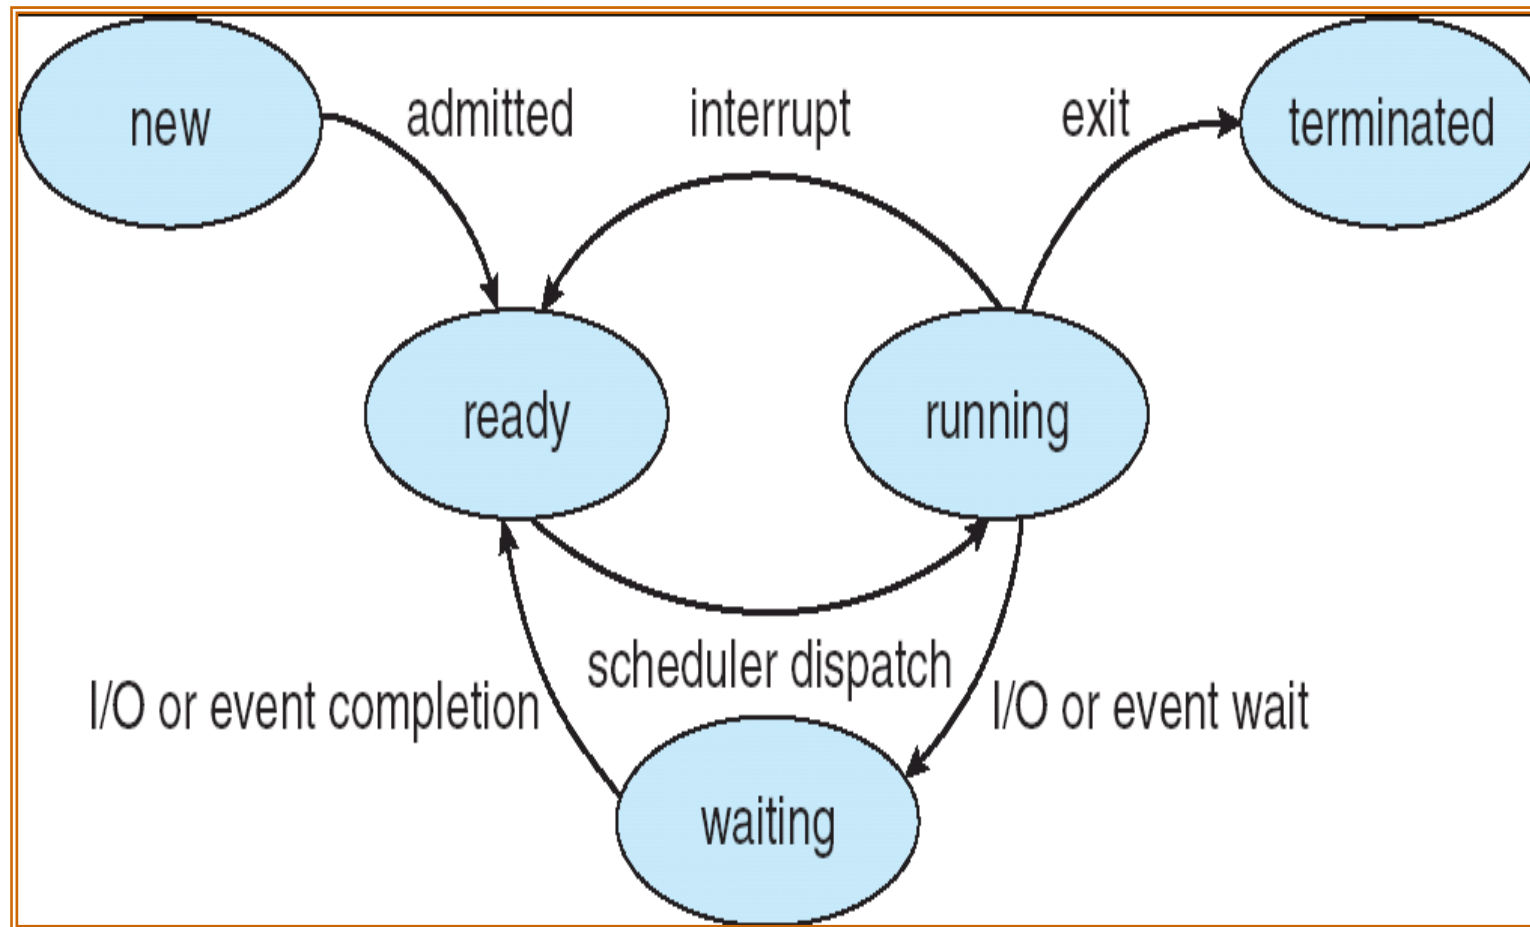
Memory Layout of a C Program



Process State

- As a process executes, it changes its *state*
 - **new**: The process is being created
 - **ready**: The process is waiting to be assigned to a CPU
 - The process is **runnable**
 - **running**: Instructions are being executed (i.e., **owns the CPU**)
 - **waiting**: The process is waiting for some event to occur
 - **terminated**: The process has finished execution

Diagram of Process State



Each process has its own process state diagram!!!

Process Control Block (PCB)

Information associated with each process

- Process state
- Program counter
 - The address of the *next* instruction
 - Must be saved when an interrupt occurs
- CPU registers
 - Vary in number and type, depending on the processor architecture
 - Accumulators, index registers, stack pointers, general purpose registers...
 - Must be saved when an interrupt occurs
- CPU scheduling information
 - Priority, pointers to scheduling queues, other scheduling parameters...

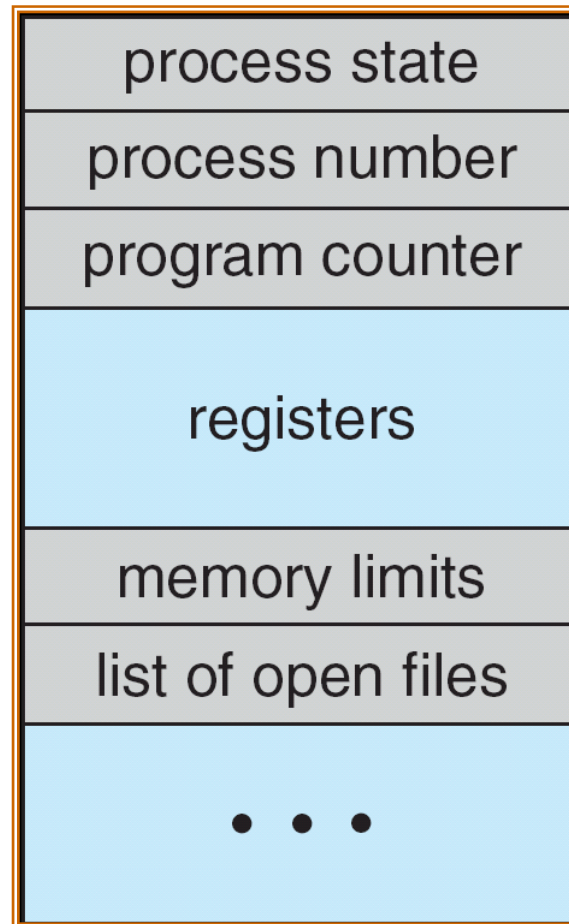
Process Control Block (PCB)

Information associated with each process (cont.)

- Memory-management information
 - Will be introduced later
- Accounting & identification information
 - CPU time and real time used
 - Time limits
 - Process numbers
- I/O status information
 - Open files and devices...

Process Control Block (PCB)

Also called **Task Control Block (TCB)**



task_struct in Linux: `/include/linux/sched.h`

Threads

- **So far**, a process is a program that performs a **single thread** of execution
 - Process executes in sequential fashion
 - Many modern operating systems support **multi-threaded** processes
 - Discussed in Chapter 4

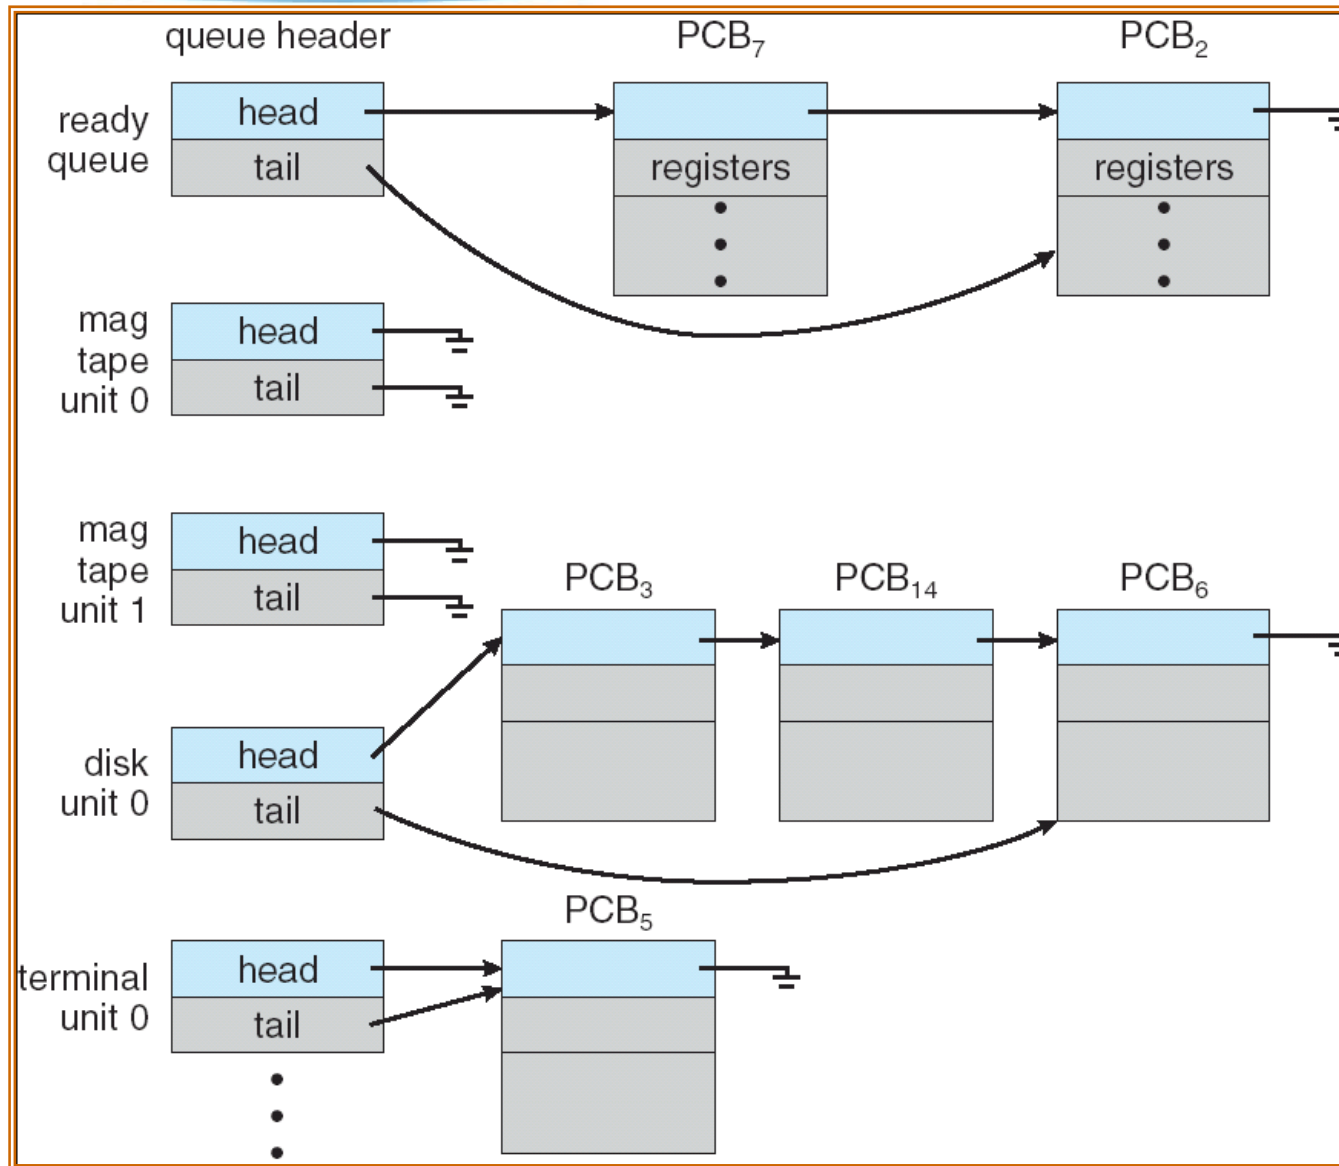
Process Scheduling

- Goals of
 - Multiprogramming
 - High CPU utilization
 - Time-sharing
 - Short response time
- A process scheduler is responsible for selecting a process to run, trying to achieve the above goals

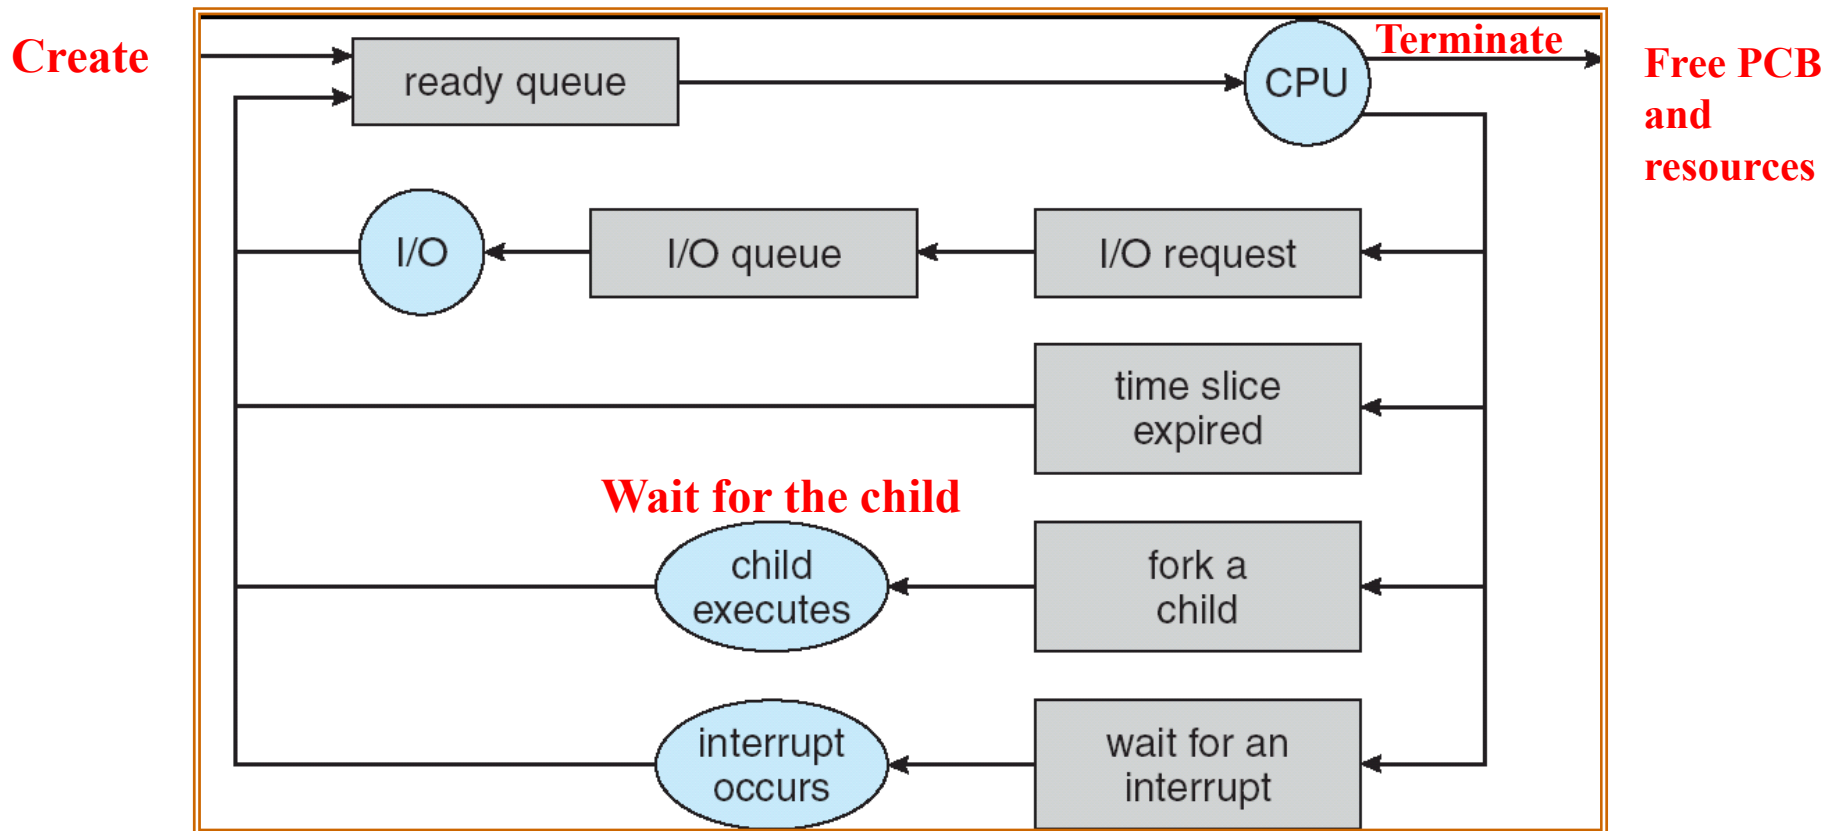
Process Scheduling Queues

- **Job queue** – set of all processes in the system
- **Ready queue (or called run queue)** – set of all processes residing in main memory, ready and waiting to execute
- **Device queues** – set of processes waiting for an I/O device
- **Event queues** – set of processes waiting for an event
- Processes migrate among the queues

Ready Queue and Various I/O Device Queues



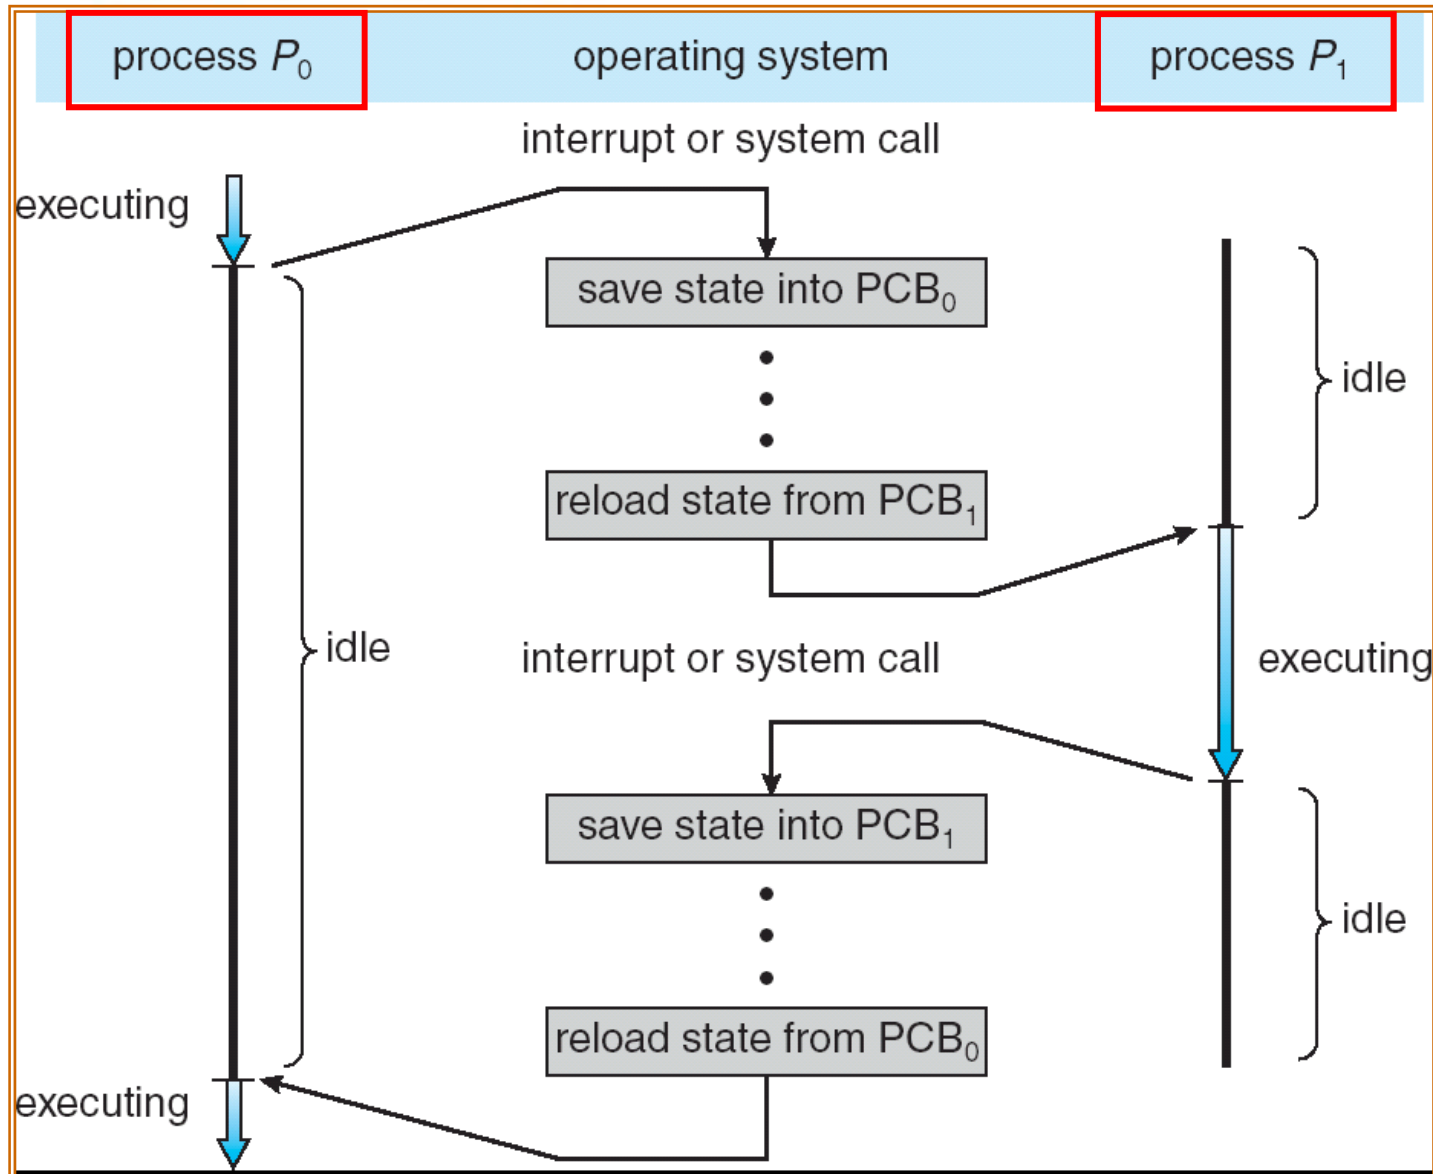
Queuing-Diagram Representation of Process Scheduling



Context Switch

- When CPU switches to another process, the system must **save** the state of the old process in its PCB and **load** the saved state for the new process
 - This is called **context switch**
- Context-switch time is overhead; the system does no useful work while switching
- Time dependent on hardware support
 - e.g., SUN UltraSPARC provide multiple set of registers
 - Context switch == change the pointer to the current register set
 - Typically, context switch requires **a few microseconds**

CPU Switch from Process to Process



Schedulers

- **Long-term scheduler** (or Job scheduler) – determines which processes should be brought into the ready queue from the job queue
- **Short-term scheduler** (or CPU scheduler) – determines which process in the ready queue should be executed next (on the CPU)

Schedulers (Cont.)

- Short-term scheduler is invoked very frequently (milliseconds) \Rightarrow (must be fast)
 - E.g., 100 ms for each process, 10ms for performing scheduling
 - $10 / (100+10) = 9\%$ CPU time wasted on the scheduling
- Long-term scheduler is invoked very infrequently (seconds, minutes) \Rightarrow (can be slow)

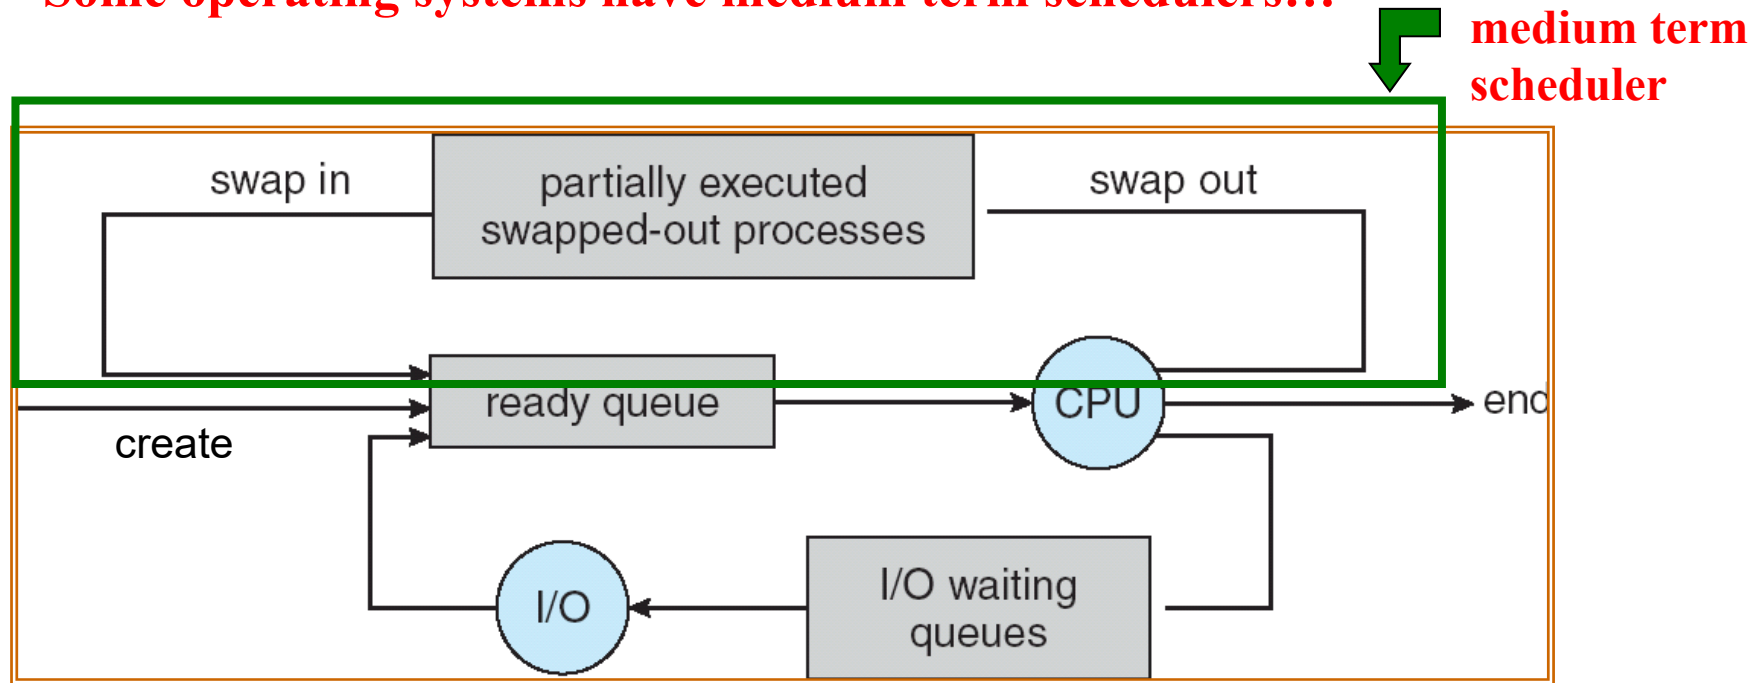
Schedulers (Cont.)



- The long-term scheduler controls the *degree of multiprogramming*
- Processes can be described as either
 - **I/O-bound process** – spends more time doing I/O than computations, many short CPU bursts
 - **CPU-bound process** – spends more time doing computations; few very long CPU bursts
- A long-term scheduler should select a good *process mix*
 - All IO-bound? All CPU-bound?
- Windows and UNIX have no long-term schedulers
 - Put all the jobs into memory

Addition of Medium Term Scheduling

Some operating systems have medium term schedulers...



Needed

for improving the process mix

when a change of memory requirement has overcommitted the available memory

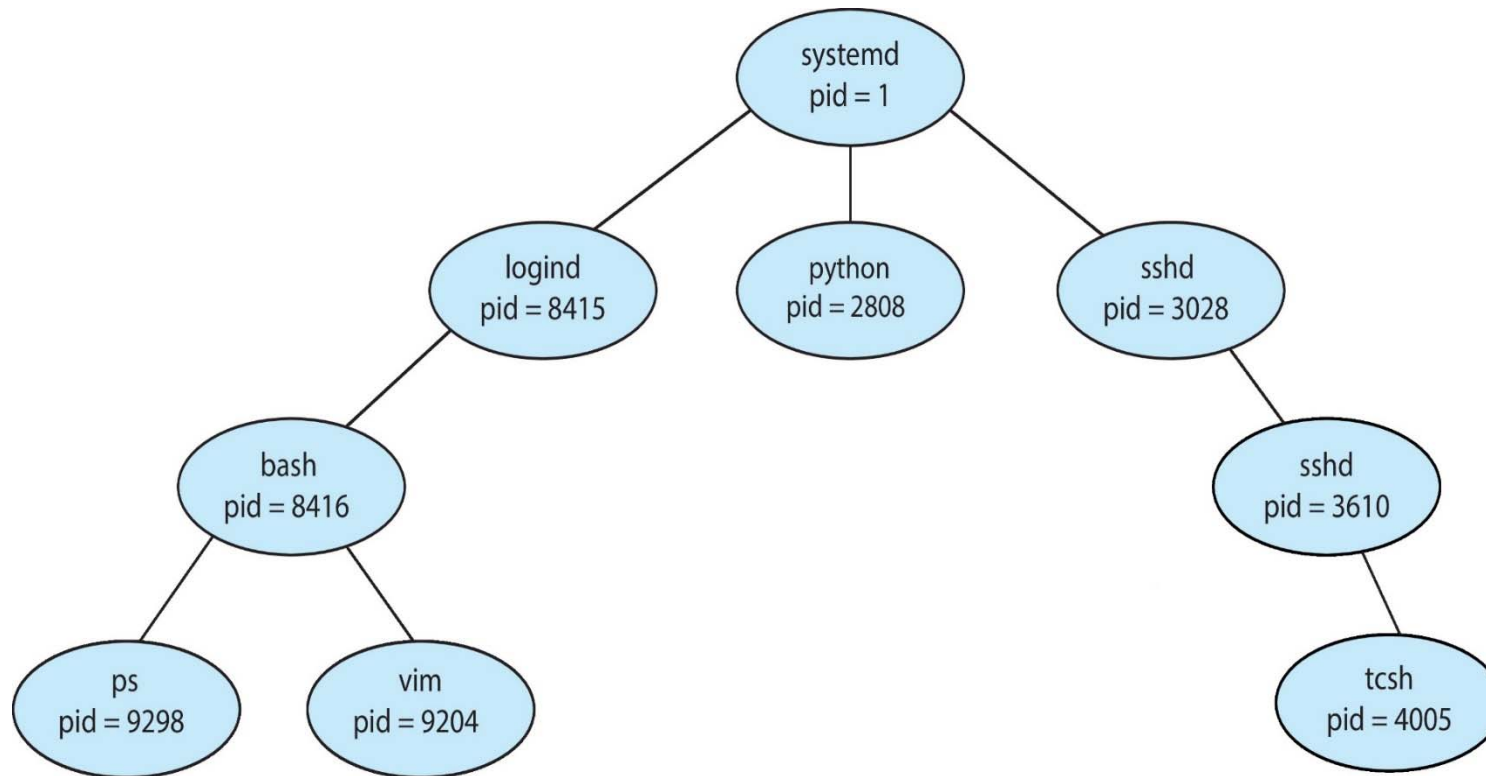
Operations on Processes



- We introduce two operations here
 - Creation
 - Termination

Process Creation

- Parent process create children processes, which, in turn create other processes, forming a **tree** of processes



Generally, a process identified via a **process identifier (pid)**

Process Creation (Cont.)

- Parent and Child
 - Resource allocation
 - A child can obtain its resources (CPU time, mem, files, IO dev...) from
 - a) OS
 - b) Its parent
 - If a child can **only** obtain its resources from its parent
 - Prevent any process from overloading the system by creating too many processes

Process Creation (Cont.)

- Parent and Child
 - Resource sharing
 - Parent and child share all resources
 - Child shares subset of parent's resources
 - Parent and child share no resources
 - Execution
 - Parent and child execute concurrently, or
 - Parent waits until child terminates
 - Address space
 - Child runs in the same address space as its parent, or
 - Child has its own address space

Process Creation (Cont.)


- UNIX examples
 - **fork** system call creates new process
 - **exec** system call
 - **replace** the process' memory space with a new program
 - usually used after a **fork**

C Program Forking Another Process

```
int main()
{
    pid_t ret, dead;
    int status;

    ret = fork();    /* fork another process */
    if (ret < 0) {    /* error occurred */
        fprintf(stderr, "Fork Failed");
        exit(-1);
    }
    else if (ret == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */ // ret == child's pid
        /* parent will wait for the child to complete */
        dead = wait (&status); // dead == the pid of the child that has died
        printf ("Child Complete");
        exit(0);
    }
}
```

Execution Flow of the Program



Process Termination

- Process executes last statement and asks the operating system to delete it (via the **exit()** system call)
 - Output data/status to its parent
 - The parent can get the exit status via the **wait()** system call
 - Process' resources are deallocated by operating system
- Parent may terminate execution of children processes
 - Reasons
 - Child has exceeded allocated resources
 - Task assigned to child is no longer required
- If parent is exiting
 - Some operating system (such as VMS) do not allow child to continue if its parent terminates
 - All children terminated - *cascading termination*
 - In UNIX, cascading termination is not required
 - Child's **parent** is set to the **init/systemd** process (pid = 1)

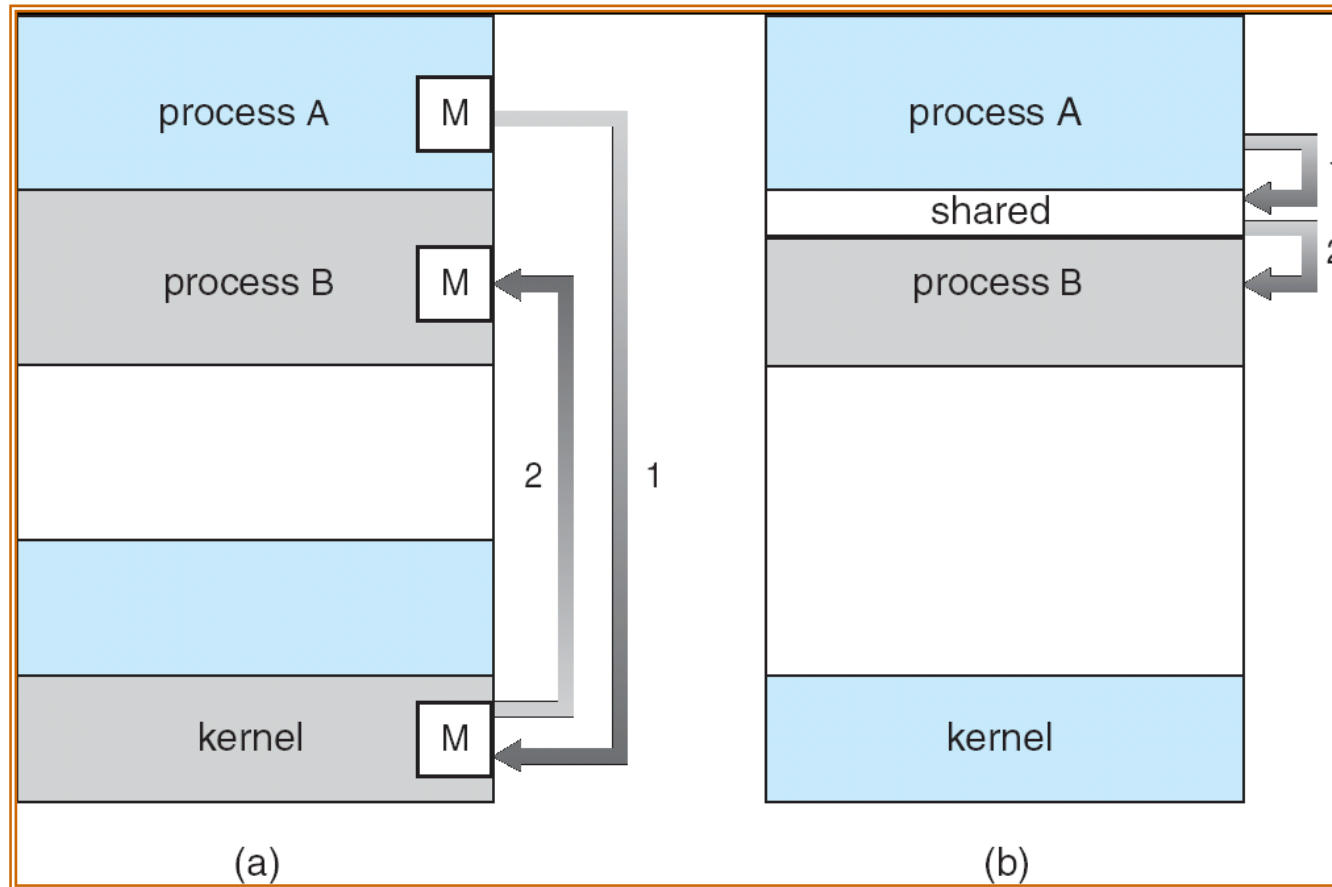
Inter-Process Communication (IPC)

- **Independent** process cannot affect or be affected by the execution of another process
- **Cooperating** process can affect or be affected by the execution of another process
 - Any processes that share data with others are cooperating processes
- Advantages of process cooperation
 - Information sharing
 - Computation speed-up
 - Divide the jobs and run them in parallel on a MP system
 - Modularity

Inter-Process Communication (IPC)

- IPC allows processes to exchange data
- Two fundamental models
 - Shared memory
 - Faster
 - Requires **no** system calls when reading/writing data
 - Message passing
 - Useful for exchanging smaller data
 - Easier to implement
 - Slower; requires system calls for sending/receiving messages

Communications Models



Message passing

Shared memory

Shared Memory

- A shared memory region must be **created** first
- Then, it is **attached** to the **address spaces** of the processes
- Memory **access** can be done freely in the region
 - Originally, OS prevents the address space of a process to be accessed by the other processes.
 - However, the users/processes of the shared memory remove this restriction
 - OS do not care about the data written in the shared memory region
- The region can be **detached** if it is no longer needed by a process

Producer-Consumer Problem

- Paradigm for cooperating processes, *producer* process produces information (i.e. data items) that is consumed by a *consumer* process
 - Use a buffer to contain the data items
 - ***unbounded-buffer*** places no practical limit on the size of the buffer
 - Producer never has to wait; consumer may need to wait
 - ***bounded-buffer*** assumes that there is a fixed buffer size
 - Both producer and consumer **may** need to wait

Bounded-Buffer – Shared-Memory Solution

- Shared data

```
#define BUFFER_SIZE 10
typedef struct {
    ...
} item;

item buffer[BUFFER_SIZE];
int in = 0;
int out = 0;
// in == out : empty, (in+1) mod BUFFER_SIZE == out: full
```

- **Solution is correct, but can only use
BUFFER_SIZE-1 elements**

Bounded-Buffer – Producer

```
while (true)
{
    ...Produce an item...
    while ( ((in + 1) % BUFFER SIZE) == out )
        ; /* do nothing -- no free buffers */
    buffer[in] = item;
    in = (in + 1) % BUFFER SIZE;
}
```

Bounded Buffer – Consumer

```
while (true)
{
    while (in == out)
        ; // do nothing -- nothing to consume

    // remove an item from the buffer
    item = buffer[out];
    out = (out + 1) % BUFFER SIZE;
    ...Process the item...
}
```

Shared Memory APIs



- System V API
 - shmget()
 - shmat()
 - shmdt()
 - shmctl()
- POSIX (Portable Operating System Interface) Shared Memory API
 - shm_open(), mmap(), munmap(), close(), shm_unlink()...

POSIX Shared Memory Example

- POSIX Shared Memory
 - Process first creates shared memory segment
shm_fd = shm_open(name, O_CREAT | O_RDWR, 0666);
 - Also used to open an existing segment
 - Set the size of the object
ftruncate(shm_fd, 4096);
 - Use **mmap()** to memory-map the shared memory object to the address space
 - Reading and writing to shared memory is done by using the pointer returned by **mmap()**

POSIX Shared Memory Example – Writing a String

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>

int main()
{
    /* the size (in bytes) of shared memory object */
    const int SIZE = 4096;
    /* name of the shared memory object */
    const char *name = "OS";
    /* strings written to shared memory */
    const char *message_0 = "Hello";
    const char *message_1 = "World!";

    /* shared memory file descriptor */
    int shm_fd;
    /* pointer to shared memory object */
    void *ptr;

    /* create the shared memory object */
    shm_fd = shm_open(name, O_CREAT | O_RDWR, 0666);

    /* configure the size of the shared memory object */
    ftruncate(shm_fd, SIZE);

    /* memory map the shared memory object */
    ptr = mmap(0, SIZE, PROT_WRITE, MAP_SHARED, shm_fd, 0);

    /* write to the shared memory object */
    sprintf(ptr,"%s",message_0);
    ptr += strlen(message_0);
    sprintf(ptr,"%s",message_1);
    ptr += strlen(message_1);

    return 0;
}
```


POSIX Shared Memory Example – Reading a String

```
#include <stdio.h>
#include <stdlib.h>
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>

int main()
{
    /* the size (in bytes) of shared memory object */
    const int SIZE = 4096;
    /* name of the shared memory object */
    const char *name = "OS";
    /* shared memory file descriptor */
    int shm_fd;
    /* pointer to shared memory object */
    void *ptr;

    /* open the shared memory object */
    shm_fd = shm_open(name, O_RDONLY, 0666);

    /* memory map the shared memory object */
    ptr = mmap(0, SIZE, PROT_READ, MAP_SHARED, shm_fd, 0);

    /* read from the shared memory object */
    printf("%s", (char *)ptr);

    /* remove the shared memory object */
    shm_unlink(name);

    return 0;
}
```

Message Passing

- Message system – processes communicate with each other without resorting to shared variables/memory
- Two major operations
 - **send**(*message*)
 - **receive**(*message*)
- Message size can be fixed or variable
- If P and Q wish to communicate, they need to
 - have a *communication link* between them
 - exchange messages via send/receive

Implementation Questions

- How are links established?
- Can a link be associated with more than two processes?
- How many links can there be between every pair of communicating processes?
- What is the capacity of a link?
- Is the size of a message fixed or variable?
- Is a link unidirectional or bi-directional?

Direct Communication

- Processes must name each other **explicitly**:
 - **send** (*P*, *message*) – send a message to process P
 - **receive**(*Q*, *message*) – receive a message from process Q
- Properties of communication link
 - Links are established **automatically**
 - Need to know each other's identity
 - A link is associated with exactly one pair of communicating processes
 - Between each pair there exists exactly one link
- The API above is **symmetric** in addressing
- Asymmetric addressing
 - **send** (*P*, *message*) – send a message to process P
 - **receive**(&*id* , *message*) – receive a message, the id **returns** the sender
- Drawback
 - Hard-coding process id
 - Changing process id would cause problems...

Indirect Communication

- Messages are sent to and received from **mailboxes** (also referred to as **ports**)
 - Each mailbox has a unique id
 - Processes can communicate only if they share a mailbox
- Properties of communication link
 - Link established only if processes share a common mailbox
 - A link may be associated with many processes
 - Each pair of processes may share several communication links
 - Link may be unidirectional or bi-directional

Indirect Communication

- Operations
 - create a new mailbox
 - send and receive messages through mailbox
 - destroy a mailbox
- Primitives are defined as:
 - send**(*A*, *message*) – send a message to **mailbox A**
 - receive**(*A*, *message*) – receive a message from **mailbox A**

Mailbox Ownership

- A mailbox is owned by the **creator** or the **OS**
 - Mailboxes owned by the Creator
 - Mailbox is part of the address space of the creator
 - The creator is the only receiver
 - Creator terminates → mailbox disappears...
 - Mailboxes owned by the OS
 - May be still existed after the creator has terminated
 - OS should provide a mechanism for users to delete the mailbox **explicitly**

Indirect Communication

- Mailbox sharing
 - P_1 , P_2 , and P_3 share mailbox A
 - P_1 sends; P_2 and P_3 receive
 - Who gets the message?
- Solutions
 - Allow a link to be associated with at most two processes
 - Allow only one process at a time to execute a receive operation
 - Allow the system to select arbitrarily the receiver. Sender is notified who the receiver was.

Synchronization

- Message passing may be either **blocking** or **non-blocking**
- **Blocking** is considered **synchronous**
 - **Blocking send** has the sender block until the message is received
 - **Blocking receive** has the receiver block until a message is available
- **Non-blocking** is considered **asynchronous**
 - **Non-blocking send** has the sender send the message and continue → **does not wait for the receiver**
 - **Non-blocking receive** has the receiver receive a valid message or null

Message Buffering

- A queue of messages attached to the link
- Implemented in one of three ways
 - Zero capacity – 0 message
Sender must wait for receiver
 - No buffering
 - Bounded capacity – finite length of n messages
Sender must wait if link full
 - Automatic buffering
 - Unbounded capacity – infinite length
Sender never waits
 - Automatic buffering
- Receivers have to wait if there is no message (in the case of blocking receive)


Message Passing APIs

- System V API
 - msgget()
 - msgsnd()
 - msgrcv()
 - msgctl()
- POSIX Message Queue API
 - mq_xxx() functions
 - mq_open(), mq_send(), mq_receive(), mq_close(), mq_unlink()....

Pipes

- A container for byte stream
- Producer writes to one end (the **write-end** of the pipe)
- Consumer reads from the other end (the **read-end** of the pipe)
- Pipes are **typically unidirectional**
 - According to POSIX.1-2001
 - On some systems, pipes are bidirectional
- Require **parent-child relationship** between communicating processes
- Windows calls these **anonymous pipes**

IPC Using Pipes



Source: <http://www.cs.columbia.edu/~jae/4118/L05-ipc.html>

IPC Using Pipes – An Example

```
...
int main(int argc, char *argv[])
{
    int pipefd[2];
    pid_t cpid;
    char buf;
    if (argc != 2) {
        fprintf(stderr, "Usage: %s <string>\n", argv[0]);
        exit(EXIT_FAILURE);
    }
    if (pipe(pipefd) == -1) {
        perror("pipe");
        exit(EXIT_FAILURE);
    }
    cpid = fork();
    if (cpid == -1) {
        perror("fork");
        exit(EXIT_FAILURE);
    }
}
```

IPC Using Pipes – An Example

```
if (cpid == 0) { /* Child reads from pipe */
    close(pipefd[1]); /* Close unused write end */
    while (read(pipefd[0], &buf, 1) > 0)
        write(STDOUT_FILENO, &buf, 1);
    write(STDOUT_FILENO, "\n", 1);
    close(pipefd[0]);
    exit(EXIT_SUCCESS);

} else { /* Parent writes argv[1] to pipe */
    close(pipefd[0]); /* Close unused read end */
    write(pipefd[1], argv[1], strlen(argv[1]));
    close(pipefd[1]); /* Reader will see EOF */
    wait(NULL); /* Wait for child */
    exit(EXIT_SUCCESS);
}

}
```

Named Pipes

- Also called FIFOs
 - On Linux, name pipes are created by **mkfifo()**
- Named pipes are more powerful than ordinary pipes
 - **No parent-child relationship** is necessary between the communicating processes
 - Any process can open the FIFO for reading or writing since it has a name
- Provided on both UNIX and Windows systems

Client-Server Communication

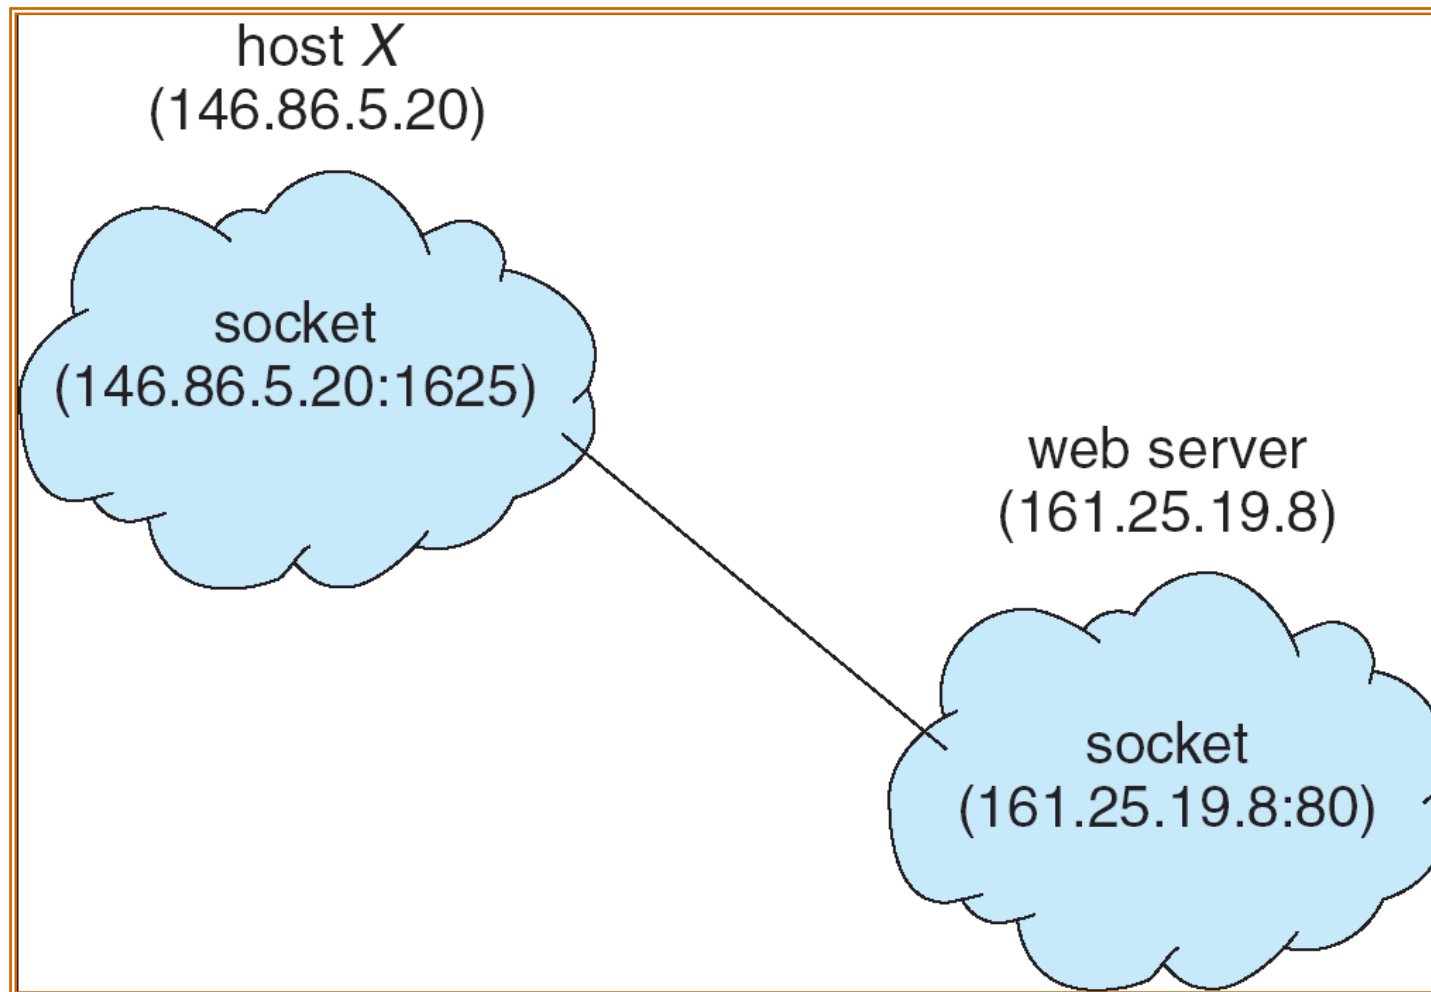


- Sockets
- Remote Procedure Calls
- Remote Method Invocation (Java)

Sockets

- A socket is defined as an *endpoint for communication*
- Concatenation of **IP address** and **port**
- The socket **161.25.19.8:1625** refers to port **1625** on host **161.25.19.8**
 - Some well known ports
 - FTP server: 21
 - HTTP server: 80
- Communication consists between a pair of sockets

Socket Communication



Remote Procedure Calls

- Remote procedure call (RPC) abstracts procedure calls between **processes on networked systems**
- Allow a client to **invoke a procedure** *on a remote host* as it would invoke a procedure locally
- **Based on message passing**
 - Each message is sent to the RPC daemon listening to a port
 - A message contains function id, and function parameters
 - Results are sent back as separated messages

Remote Procedure Calls

- **Stubs**
 - transforms procedure calls to msgs, and vice versa.
 - Hide the communication details
- The **client-side stub** locates the server and *marshals* the parameters
- The **server-side stub** receives this message, unpacks the marshaled parameters, and performs the procedure on the server
- The return values are processed in a similar manner

Issues about RPC

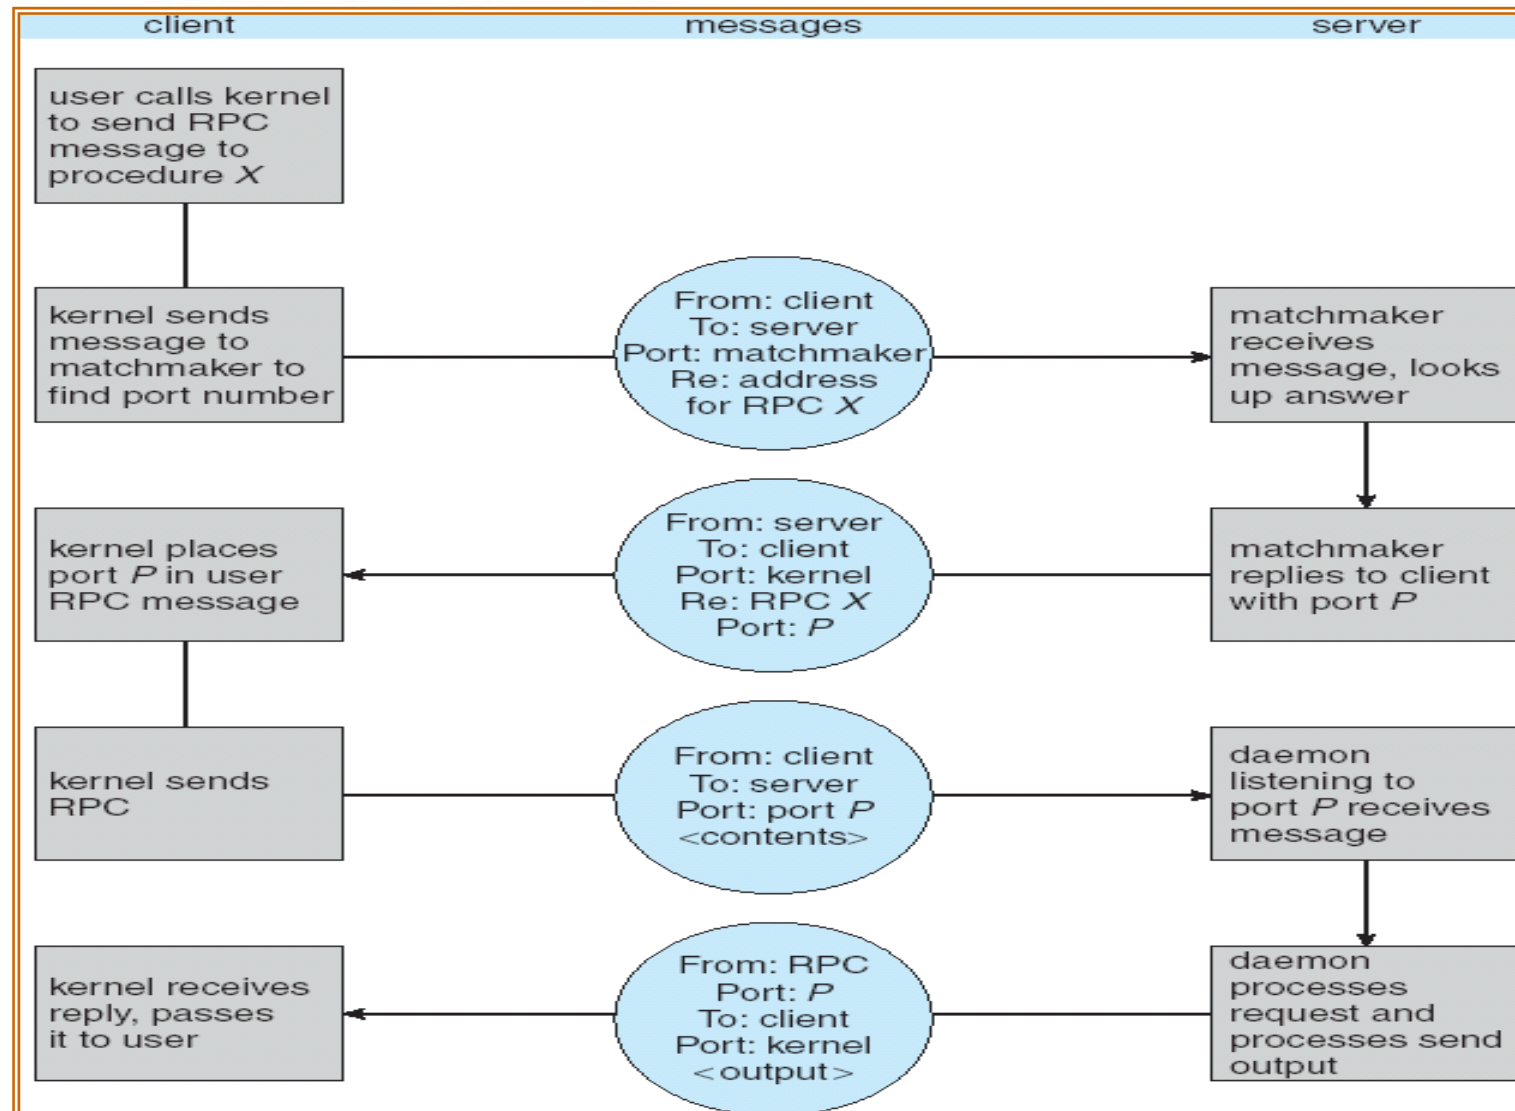


- Data representation
 - Little Endian? Big Endian?
 - External Data Representation (XDR)
- Error handling
 - RPC can **fail**, or be **executed more than once** due to **network error**
 - OS ensures the “exactly once” semantic by (ALO && AMO)
 - At least once (ALO)
 - Server **acks** the client when the client msg has been received & executed
 - Client msg transmission repeats until ACKed
 - At most once (AMO)
 - Each msg has a timestamp
 - Drop the msg if its timestamp is in the **history** on the server

Issues about RPC

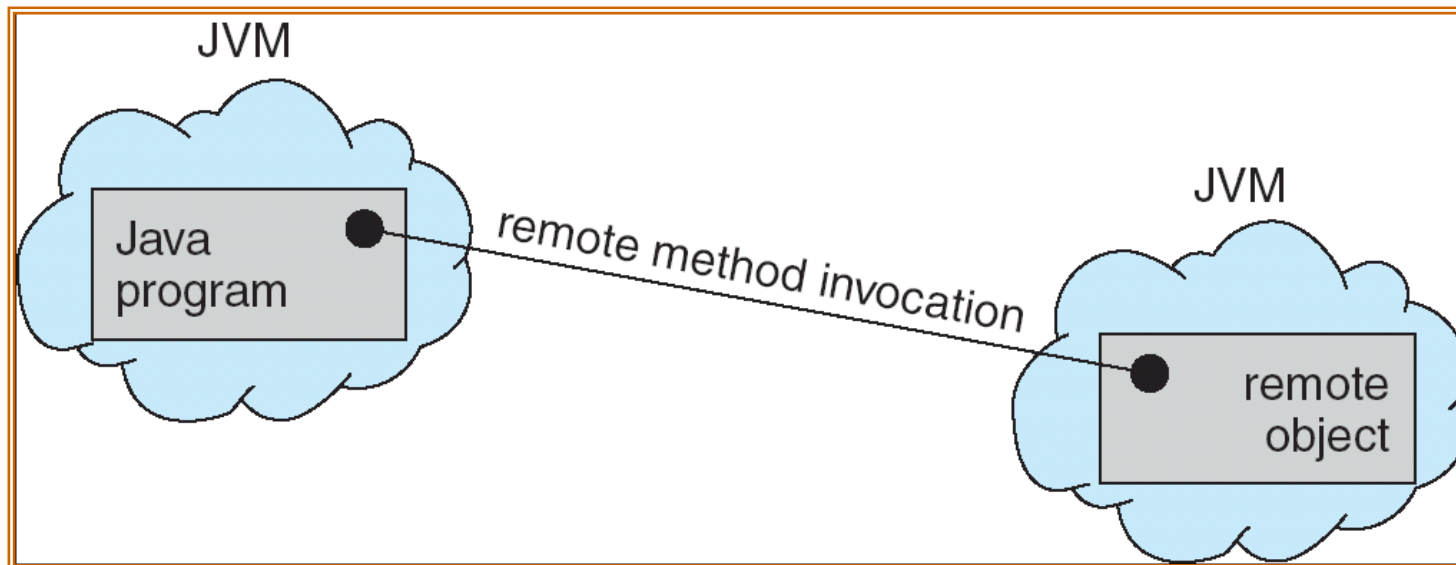
- Binding
 - How does the client know the port number of the server?
 - Fixed port number
 - Inflexible
 - » Server can not change its port easily
 - Dynamic binding
 - A rendezvous daemon (a [matchmaker](#)) listening on a fixed port to find out the port of the server
 - See the next slide

Execution of RPC



Remote Method Invocation

- Remote Method Invocation (RMI) is a Java mechanism similar to RPCs
- RMI allows a Java program on one machine to invoke a method on a remote object
- Different from RPC, RMI is
 - Object-based
 - Possible to pass objects as parameters



Marshalling Parameters

