

# Chapter 8

# Deadlocks

**Da-Wei Chang**

**CSIE.NCKU**

*Source: Abraham Silberschatz, Peter B. Galvin, and Greg Gagne,  
"Operating System Concepts", 10th Edition, Wiley.*

# Outline



- System Model
- Deadlock Characterization
- Methods for Handling Deadlocks
- Deadlock Prevention
- Deadlock Avoidance
- Deadlock Detection
- Recovery from Deadlock

# Chapter Objectives

- To develop a description of deadlocks, which prevent **sets of concurrent processes** from completing their tasks
- To present a number of different methods for **preventing** or **avoiding** deadlocks in a computer system

# System Model

- Resource types  $R_1, R_2, \dots, R_m$   
*CPU cycles, memory space, I/O devices*
- Each resource type  $R_i$  has  $W_i$  instances
- Each process utilizes a resource as follows
  - request
  - use
  - release
- A deadlock example
  - P1 uses R1, requests R2
  - P2 uses R2, requests R1

# Deadlock Characterization

Deadlock can occur if four conditions hold **simultaneously**.

- **Mutual exclusion**

- only one process at a time can use a resource

- **Hold and wait**

- a process holding at least one resource is waiting to acquire additional resources held by other processes

- **No preemption**

- a resource can be released only **voluntarily** by the process holding it

- **Circular wait**

- there exists a set of waiting processes  $\{P_1, P_2, \dots, P_n\}$  such that  $P_1$  is waiting for a resource held by  $P_2$ ,  $P_2$  is waiting for a resource held by  $P_3$ , ...,  $P_{n-1}$  is waiting for a resource held by  $P_n$ , and  $P_n$  is waiting for a resource held by  $P_1$


# Resource-Allocation Graph

A set of vertices  $V$  and a set of edges  $E$

- $V$  is partitioned into two types
  - $P = \{P_1, P_2, \dots, P_n\}$ , the set consisting of all the processes in the system
  - $R = \{R_1, R_2, \dots, R_m\}$ , the set consisting of all resource types in the system
- *request* edge – directed edge  $P_i \rightarrow R_j$
- *assignment* edge – directed edge  $R_j \rightarrow P_i$

# Resource-Allocation Graph (Cont.)


- Process




- Resource Type with 4 instances




- $P_i$  requests an instance of  $R_j$



- $P_i$  is holding an instance of  $R_j$




# Example of a Resource Allocation Graph






# Resource Allocation Graph with a Deadlock



# Resource Allocation Graph with a Cycle **But No Deadlock**



# Basic Facts

- If graph contains no cycles  $\Rightarrow$  no deadlock
- If graph contains a cycle  $\Rightarrow$ 
  - if only one instance per resource type, then deadlock
  - if several instances per resource type, possibility of deadlock

# Methods for Handling Deadlocks

- Ensure that the system will *never* enter a deadlock state
  - Deadlock *prevention* or *avoidance*
- Allow the system to enter a deadlock state and then *recover*
- *Ignore* the problem and pretend that deadlocks never occur in the system
  - used by most operating systems (UNIX, Windows..)
  - applications have to deal with deadlocks themselves

# Deadlock Prevention

- **Restrain the ways request can be made**
  - Ensure at least one of the four conditions cannot hold
- **Mutual Exclusion**
  - not required for sharable resources; must hold for non-sharable resources
- **Hold and Wait** – must guarantee that whenever a process requests a resource, it does not hold any other resources
  - Require process to request and **be allocated all its resources before it begins execution**, or allow process to **request resources only when the process has none**
  - Low resource utilization
  - Starvation possible
    - A process that needs several popular resources may wait infinitely

# Deadlock Prevention (Cont.)



- **No Preemption –**

- If a process **holds some resources** and **requests another resource that cannot be immediately allocated** to it, then all resources currently being held are **released** (i.e., **preempt** the resource holding)
  - Preempted resources are added to the list of resources for which the process is waiting
  - State rollback may be required
  - Process will be **restarted only when** it can **regain its old resources, as well as the new one** that it is requesting

- **Circular Wait –**

- impose a **total ordering** of all resource types, and require that each process **requests resources in an increasing order** of enumeration
  - Lock-order verifier. example: witness (BSD)

# Deadlock Avoidance

- Check if the **admission of the current request** may lead to a deadlock state (specifically, a circular-wait condition)
  - Yes → do not admit the request
  - Requires that the system has some additional *a priori information* available
  - Simplest and most useful model requires
    - each process declare the *maximum number of resources* of each type that it may need
  - Resource-allocation *state* is defined by the number of **available** and **allocated** resources, and the **maximum demands** of the processes

# Safe State

- When a process requests an available resource, system must decide if the immediate allocation leaves the system in a **safe state**
- System is in safe state if there exists a **safe sequence** of all processes
  - Sequence  $\langle P_1, P_2, \dots, P_n \rangle$  is **safe** if for each  $P_i$ , the resources that  $P_i$  can still request can be satisfied by currently available resources + resources held by all the  $P_j$ , where  $j < i$ .
    - If  $P_i$ 's resource needs are not immediately available, then  $P_i$  can wait until all  $P_j$  have finished.
    - When  $P_j$  is finished,  $P_i$  can obtain needed resources, execute, return allocated resources, and terminate.
    - When  $P_i$  terminates,  $P_{i+1}$  can obtain its needed resources, and so on.




# Basic Facts



- If a system is in safe state  $\Rightarrow$  no deadlocks
- If a system is in unsafe state  $\Rightarrow$  possibility of deadlock
- Avoidance  $\Rightarrow$  ensure that a system will never enter an unsafe state

# Safe, Unsafe, and Deadlock States




# Deadlock Avoidance Algorithms

- Resource-Allocation Graph Algorithm
- Banker's Algorithm

# Resource-Allocation Graph Algorithm


- *Claim edge*  $P_i \rightarrow R_j$  indicated that process  $P_i$  **MAY** request resource  $R_j$ 
  - represented by a **dashed line**
- Claim edge converted to **request edge** when a process requests a resource
- When a resource is released by a process, the assignment edge is reconverted to a claim edge
- Resources must be claimed *a priori* in the system
- A request is granted **if it does not cause a cycle**
  - Safe state

# Resource-Allocation Graph for Deadlock Avoidance



**Safe**

# Unsafe State in Resource-Allocation Graph



**P2 request R2**

- ➔ Granting the request will cause a cycle
- ➔ the request cannot be granted

\*This method does not support resources with multiple instances

# Banker's Algorithm

- Supports resources with multiple instances
- Each process must *a priori* claim maximum use
- When a process requests a resource it may have to wait

# Data Structures for the Banker's Algorithm

Let  $n$  = number of processes, and  $m$  = number of resources types.

- *Available*: Vector of length  $m$ . If available  $[j] = k$ , there are  $k$  instances of resource type  $R_j$  available.
- *Max*:  $n \times m$  matrix. If  $Max[i,j] = k$ , then process  $P_i$  may request at most  $k$  instances of resource type  $R_j$ .
- *Allocation*:  $n \times m$  matrix. If  $Allocation[i,j] = k$  then  $P_i$  is currently allocated  $k$  instances of  $R_j$ .
- *Need*:  $n \times m$  matrix. If  $Need[i,j] = k$ , then  $P_i$  may need  $k$  more instances of  $R_j$  to complete its task.
  - $Need[i,j] = Max[i,j] - Allocation[i,j]$



# Resource-Request Algorithm for Process $P_i$

$Request_i$  = request vector for process  $i$

$Request_i[j] = k \Rightarrow$  process  $P_i$  wants  $k$  instances of resource type  $R_j$

1. If  $Request_i \leq Need_i$ , go to step 2. Otherwise, raise error condition, since process has exceeded its maximum claim.
2. If  $Request_i \leq Available$ , go to step 3. Otherwise,  $P_i$  must wait, since resources are not available.
3. **Pretend** to allocate requested resources to  $P_i$  by modifying the state as follows:

$$Available = Available - Request_i$$

$$Allocation_i = Allocation_i + Request_i$$

$$Need_i = Need_i - Request_i$$

4. Invoke the safety algorithm
  - If safe  $\Rightarrow$  the resources are allocated to  $P_i$
  - If unsafe  $\Rightarrow P_i$  must wait, and the old resource-allocation state is restored

# Safety Algorithm

1. Let *Work* and *Finish* be vectors of length  $m$  and  $n$ , respectively

Initialization--

$Work = Available$

$Finish[i] = false$  for  $i = 0, 1, 2, \dots, n$

2. Find an  $i$  such that both

(a)  $Finish[i] = false$

(b)  $Need_i \leq Work$  // can allow it to finish

If no such  $i$  exists, go to step 4

3.  $Work = Work + Allocation_i$  //  $P_i$  releases its resources

$Finish[i] = true$

go to step 2

4. If  $Finish[i] == true$  for all  $i$ , then the system is in a safe state

# Example of Banker's Algorithm

- 5 processes  $P_0$  through  $P_4$
- 3 resource types
  - $A$  (10 instances),  $B$  (5 instances), and  $C$  (7 instances).
- Snapshot at time  $T_0$

	<u>Allocation</u>	<u>Max</u>	<u>Available</u>
	$A \ B \ C$	$A \ B \ C$	$A \ B \ C$
$P_0$	0 1 0	7 5 3	3 3 2
$P_1$	2 0 0	3 2 2	
$P_2$	3 0 2	9 0 2	
$P_3$	2 1 1	2 2 2	
$P_4$	0 0 2	4 3 3	

# Example (Cont.)

- The content of the matrix. **Need** is defined to be Max - Allocation

	<u>Need</u>		
	<i>A</i>	<i>B</i>	<i>C</i>
$P_0$	7	4	3
$P_1$	1	2	2
$P_2$	6	0	0
$P_3$	0	1	1
$P_4$	4	3	1

- The system is **in a safe state** since the sequence  $\langle P_1, P_3, P_4, P_2, P_0 \rangle$  satisfies safety criteria

# Example $P_1$ Request (1,0,2) (Cont.)

- Check that Request  $\leq$  Available (that is,  $(1,0,2) \leq (3,3,2)$ )  $\Rightarrow$  true

	<u>Allocation</u>	<u>Need</u>	<u>Available</u>
	A B C	A B C	A B C
$P_0$	0 1 0	7 4 3	2 3 0
$P_1$	3 0 2	0 2 0	
$P_2$	3 0 2	6 0 0	
$P_3$	2 1 1	0 1 1	
$P_4$	0 0 2	4 3 1	

- Executing safety algorithm shows that sequence  $\langle P_1, P_3, P_4, P_0, P_2 \rangle$  satisfies safety requirement
- Then,
  - can request for (3,3,0) by  $P_4$  be granted? No,  $>$  available
  - can request for (0,2,0) by  $P_0$  be granted? No, unsafe

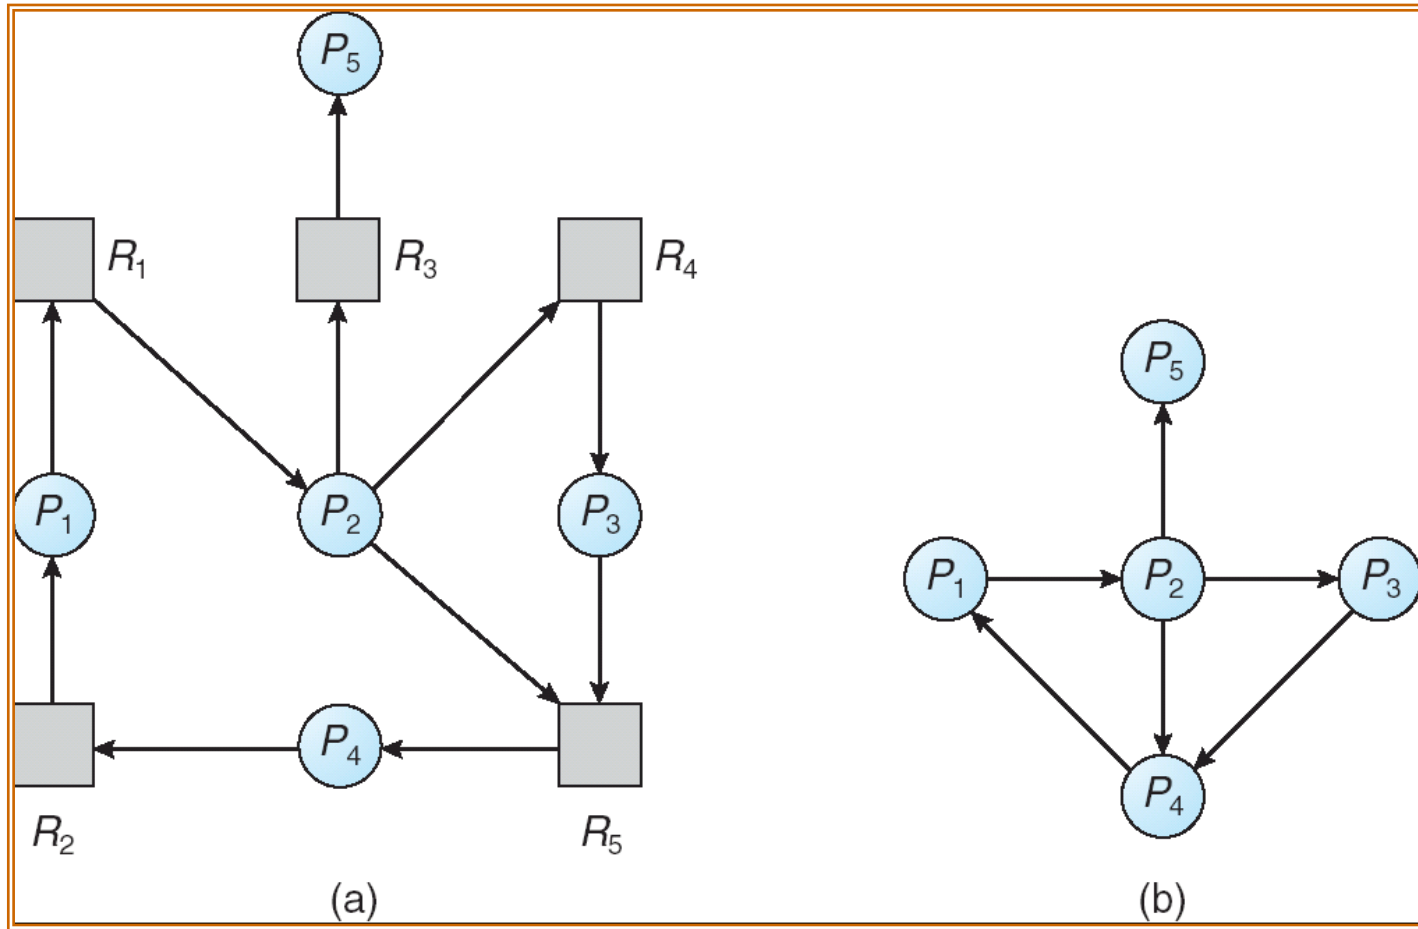
# Deadlock Detection

- Allow system to enter deadlock state
- Detection algorithm
- Recovery scheme

# Single Instance for Each Resource Type

- Maintain *wait-for* graph
  - Nodes are processes
  - $P_i \rightarrow P_j$  if  $P_i$  is waiting for  $P_j$
- Periodically invoke an algorithm that searches for a cycle in the graph
- An algorithm to detect a cycle in a graph requires an order of  $n^2$  operations, where  $n$  is the number of vertices in the graph

# Resource-Allocation Graph and Wait-for Graph



Resource-Allocation Graph

Corresponding wait-for graph



# Several Instances of a Resource Type

- *Available*: A vector of length  $m$  indicates the number of available resources of each type
- *Allocation*: An  $n \times m$  matrix defines the number of resources of each type currently allocated to each process
- *Request*: An  $n \times m$  matrix indicates the **current request** of each process. If  $Request[i, j] = k$ , then process  $P_i$  is requesting  $k$  more instances of resource type  $R_j$
- Different from the banker's algorithm used for deadlock avoidance
  - Check **current request**, instead of **future need**
  - **Not always** check when a request is made

# Detection Algorithm

1. Let *Work* and *Finish* be vectors of length  $m$  and  $n$ , respectively.

Initial condition:

(a) *Work* = *Available*

(b) For  $i = 1, 2, \dots, n$ ,

    If  $Allocation_i \neq 0$ ,  $Finish[i] = \text{false}$

    Otherwise,  $Finish[i] = \text{true}$

2. Find an index  $i$  such that both:

(a)  $Finish[i] == \text{false}$

(b)  $Request_i \leq Work$

If no such  $i$  exists, go to step 4.

# Detection Algorithm (Cont.)

3.  $Work = Work + Allocation_i$   
 $Finish[i] = true$   
go to step 2.
4. If  $Finish[i] == false$ , for some  $i$ ,  $1 \leq i \leq n$ , then the system is in deadlock state. Moreover, if  $Finish[i] == false$ , then  $P_i$  is in a deadlock state.

Optimistically assumes that  $P_i$  will release its resources after the request.

If the assumption is not correct, deadlock may occur later.

- detected by the later invocation of the detection algorithm

Algorithm requires an order of  $O(m \times n^2)$  operations to detect whether the system is in a deadlock state.

# Example of Detection Algorithm

- Five processes  $P_0$  through  $P_4$ ; three resource types A (7 instances), B (2 instances), and C (6 instances).
- Snapshot at time  $T_0$ :

	<u>Allocation</u>			<u>Request</u>			<u>Available</u>		
	A	B	C	A	B	C	A	B	C
$P_0$	0	1	0	0	0	0	0	0	0
$P_1$	2	0	0	2	0	2			
$P_2$	3	0	3	0	0	0			
$P_3$	2	1	1	1	0	0			
$P_4$	0	0	2	0	0	2			

- Sequence  $\langle P_0, P_2, P_3, P_1, P_4 \rangle$  will result in  $Finish[i] = \text{true}$  for all  $i$ .

# Example (Cont.)

- $P_2$  requests an additional instance of type  $C$ .

	<u>Request</u>		
	$A$	$B$	$C$
$P_0$	0	0	0
$P_1$	2	0	2
$P_2$	0	0	1
$P_3$	1	0	0
$P_4$	0	0	2

- Is the system in a safe state?
  - Can reclaim resources held by process  $P_0$ , but insufficient resources to fulfill other processes
  - Deadlock exists, consisting of processes  $P_1$ ,  $P_2$ ,  $P_3$ , and  $P_4$

# Detection-Algorithm Usage

- When, and how often, to invoke? It depends on
  - How often a deadlock is likely to occur?
  - How many processes will be affected?
- Possible ways
  - Invoked at each request that needs to wait
  - Invoked periodically
  - Invoked at specific situations
    - E.g., when the system throughput suddenly drops

# Recovery from Deadlock



- Two general approaches
  - Process Termination
  - Resource Preemption

# Recovery from Deadlock: Process Termination

- Abort all deadlock processes
- Abort one process at a time until the deadlock cycle is eliminated
- Selecting a victim
  - Priority of the process
  - How long process has computed, and how much longer to completion
  - Resources the process holds
  - Resources a process needs to complete



# Recovery from Deadlock: Resource Preemption

- Preempt some resources from some processes and give the resources to other processes
- Selecting a victim – minimize cost, similar to the previous slide
- Rollback – return to some *safe state*, restart process from that state
- Starvation – a process may always be selected as the victim
  - May need to include *number of rollback* in the cost factor