

#### **Outline**

- Background
- Demand Paging
- Copy-on-Write
- Page Replacement
- Allocation of Frames
- Thrashing
- Memory-Mapped Files
- Allocation Kernel Memory
- Other Considerations
- Operating System Examples

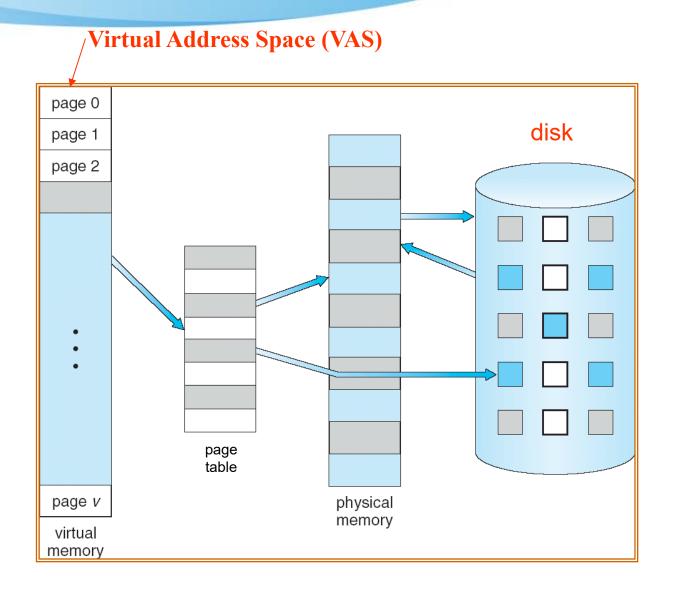
- So far, a program need to be entirely loaded into the memory before it can execute
  - Limit the number of processes in the memory
  - Process memory can not be lager than physical memory
- However, part of a program may be rarely used
  - Error code, unusual routines, large data structures
- So, the entire program does not needed at the same time

- Consider ability to execute partially-loaded program
  - Program no longer constrained by limits of physical memory
  - Each program takes less memory while running →
     more programs run at the same time
    - improve performance
  - Less I/O needed to load/swap programs into memory
    - → faster program loading, shorter program startup time

- Virtual memory separation of user logical memory from physical memory
  - Only part of the program needs to be in memory for execution
  - Logical address space can therefore be much larger than physical address space
  - Allows address spaces to be shared by several processes
    - As shown in Chapter 9
  - Allows for more efficient process creation
    - Less IO for program loading
  - Allows more processes in the memory
  - Makes the task of programming much easier
    - Programmer no longer needs to worry about the amount of physical memory available

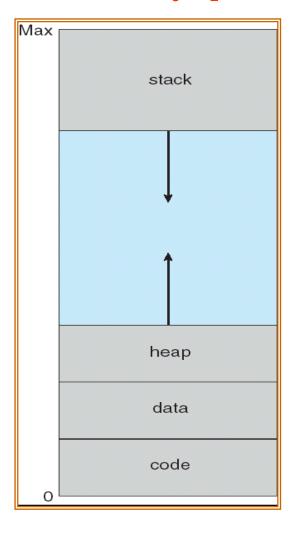
- Virtual memory can be implemented via
  - Demand paging (introduced later)
  - Demand segmentation (introduced later)

## Virtual Memory > Physical Memory



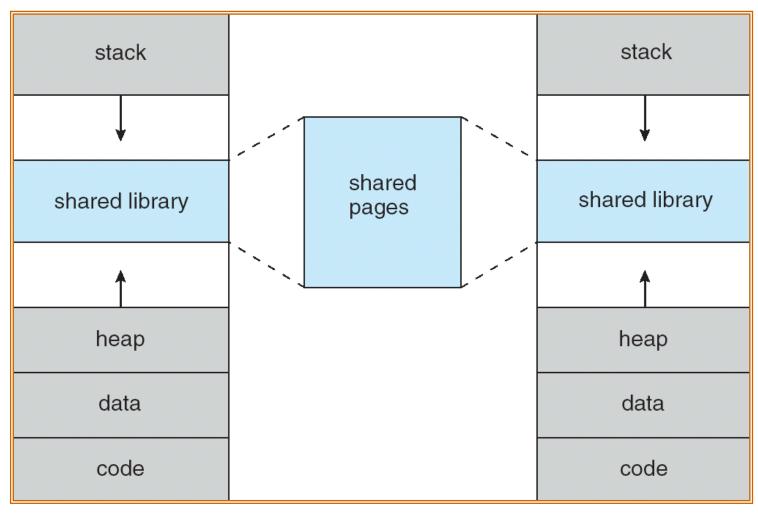
## Virtual Address Space (VAS)

#### VAS is usually sparse...



- ■Usually, stack starts at max logical address and grow "down" while heap grows "up"
  - •Unused address space between the two is hole
    - No physical memory needed for the hole
  - The hole is left for heap/stack growth, dynamically linked shared libraries, and etc.
- System libraries shared via mapping into virtual address space (see next slide)
- Shared memory by mapping pages read-write into virtual address space

# 

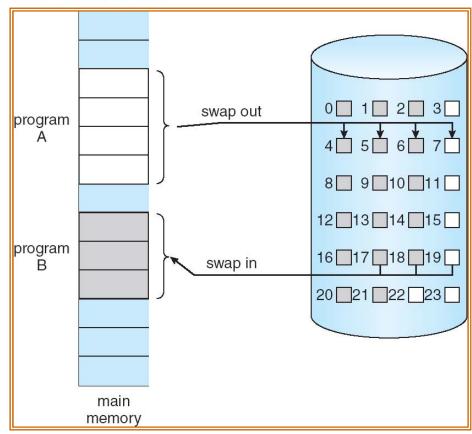


Shared memory can be implemented in a similar way

#### **Demand Paging**

- Bring a page into memory only when it is needed
  - Less I/O needed
  - Less memory needed
  - Faster response
  - More processes
- Page is needed? ⇒ reference to the page
- Page reference
  - valid, in-memory  $\Rightarrow$  access the page
  - valid, not-in-memory  $\Rightarrow$  bring to memory (Page Fault)
  - invalid reference  $\Rightarrow$  abort/exception (Seg. Fault)

# Transferring Pages between Memory and Disk

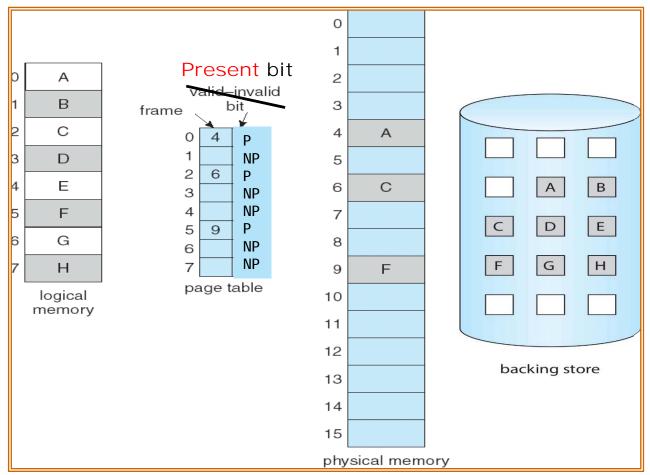


Demand paging is similar to a paging system with swapping However, it uses a *lazy* swapper → swap in a page only when the page is needed

The lazy swapper is called *pager*...

### **Not-in-Memory Pages**

#### How do we know if a page is in memory or not?



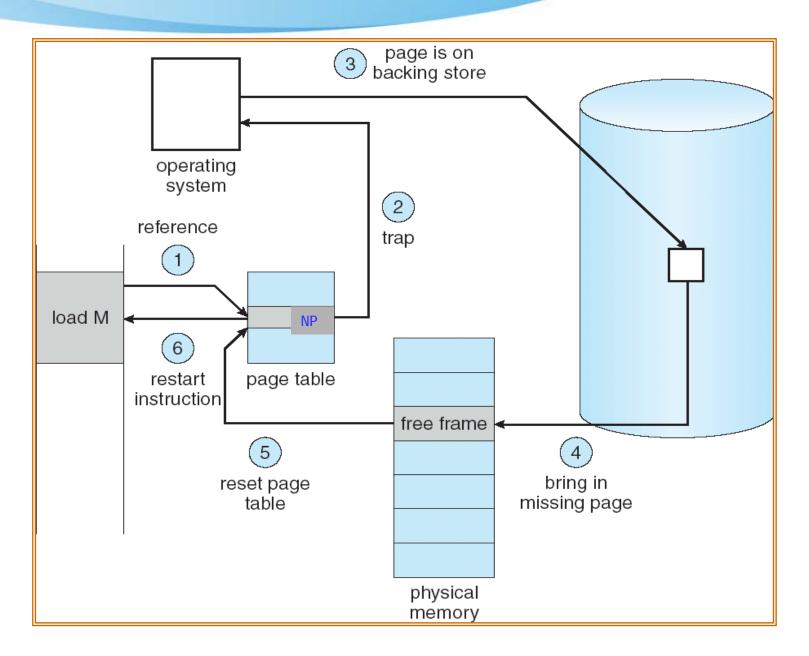
A page fault occurs when a process accesses a not-in-memory (valid) page

## Steps in Handling a Page Fault

- 1. Reference to a non-in-memory page → trap to OS
  - "Page fault"
- 2. OS checks its data structures to decide
  - a) Invalid reference  $\Rightarrow$  abort (segmentation fault), or
  - b) Just not in memory  $\Rightarrow$  page-in the page (the REAL page fault)
- 3. Find free frame
  - If no free frame, page-out a used page to make room
    - page replacement
- 4. Page-in the page into a free frame via scheduled disk operation
- 5. Update PTE and MMU to record the virtual-physical mapping
- 6. Restart the instruction that caused the page fault

Page faults are transparent to processes!

## Steps in Handling a Page Fault



#### **Detailed Page Fault Steps**

- 1. Trap to the operating system
- 2. Save the user registers and process state
- 3. Determine that the interrupt was a page fault
- 4. Check that the page reference was legal and determine the location of the page on the disk
- 5a. Read the page from the disk to a free frame
  - a) Send a read request to the disk, wait for the disk queue until the read request is serviced
  - b) Wait for the disk seek and/or latency time
  - c) Begin the transfer of the page to a free frame (DMA)
- 5b. While waiting, allocate the CPU to other process(es)
  - Rescheduling, change the *current process*

#### **Detailed Page Fault Steps**

- 6. The I/O completed, receive an interrupt (from the disk)
- 7. Save the registers and state of the current process
- 8. Determine that the interrupt was from the disk
- 9. Update the page table and TLB to record the virtual-physical mapping
- 10. Wait for the CPU to be allocated to this process again
- 11. Restore the user registers & process state, and then resume the interrupted instruction

#### Performance of Demand Paging

- Page Fault Rate  $0 \le p \le 1$ 
  - if p = 0, no page faults
  - if p = 1, every reference causes a fault
- Effective Access Time (EAT)

EAT = 
$$(1 - p)$$
 \* memory access time  
+  $p$  \* page fault time

- Page fault time
  - Service the page fault exception (short time)
  - page in/out (long time)
  - restart the process (short time)

#### **Demand Paging Example**

- Memory access time = 200 ns
- Average page-fault time = 8 ms

• EAT = 
$$(1-p) * 200 \text{ns} + p * (8 \text{ ms})$$
  
=  $(1-p) * 200 + p * 8,000,000$   
=  $200 + p * 7,999,800$ 

- If  $p = 0.001 \rightarrow EAT = 8.2 \text{ us} \rightarrow 40 \text{x slowdown!}$
- If performance degradation < 10%

$$-200 + 7,999,800 * p < 220$$
 →  $7,999,800 * p < 20$  →  $p < .0000025$ 

- < one page fault in every 400,000 memory accesses!!!</p>

#### **Aspects of Demand Paging**

- Extreme case pure demand paging
  - never bring a page into memory until it is required
  - a process is started with no pages in memory
- Code pages can be paged in from the executable file, and discarded (rather than paging out, i.e., writing back to disk) when freeing frames
  - Used in Solaris and current BSD
  - Swap space is still required for the other types of pages (e.g., data, stack, heap...)

#### Implementations of fork()

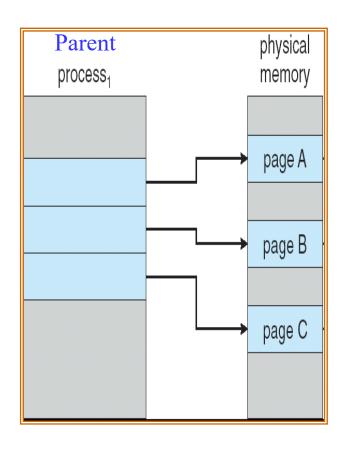
- Three types of implementations
  - Traditional fork
  - vfork
  - COW (Copy-on-Write) based fork

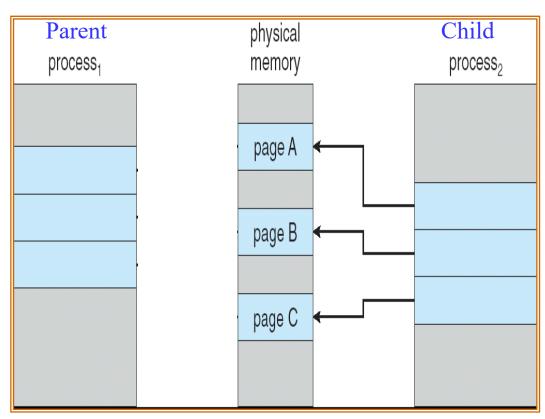
- Traditional fork
  - Copy all the frames from parent to child
  - Large memory copy overhead

#### vFork

- Intended to be used when the child calls *exec()*
- The parent process is suspended
- The child process *borrows* frames from the parent
  - Changes to the pages are visible to the parent once the parent resumes
- The borrowed frames are returned when the child call *exec()*

#### vFork



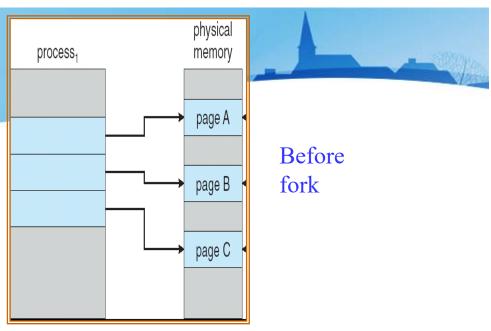


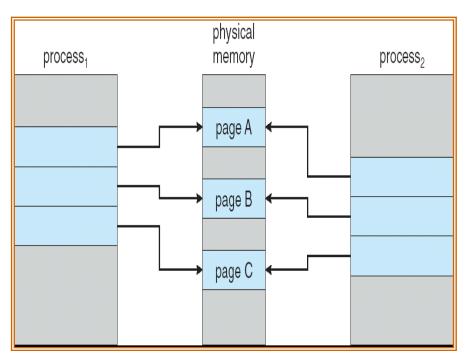
Before fork After fork

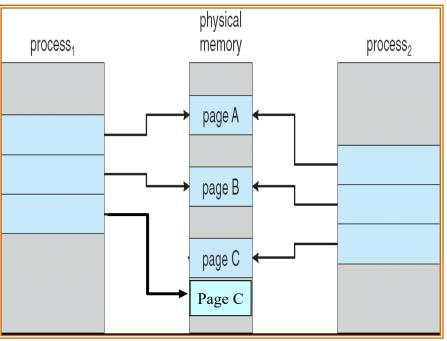
#### Copy-on-Write

- Copy-on-Write (COW) allows both parent and child processes to initially *share* the same pages in memory
  - A way to implement fork()
  - Based on page fault (but another type of page fault)
    - (X) Not-in-memory page fault
    - (V) Protection violation page fault
- The shared page is set as read-only initially. If either process modifies the shared page
  - Page fault occurs, and
  - The page is copied
- COW allows more efficient process creation as only modified pages are copied

# Copy-on-Write Fork







After fork

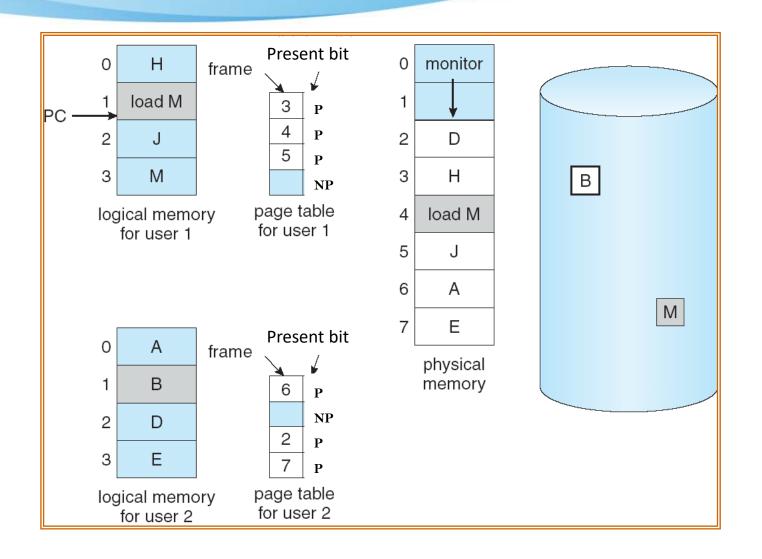
When parent tries to update page C

#### Page Replacement

• Swap in a page when all the frames are occupied → page replacement is needed

 Page-fault service routine needs to deal with page replacement


### Need for Page Replacement



#### **Basic Page Replacement**


- 1. Find the location of the desired page on disk
- 2. Find a free frame
  - If there is a free frame, use it
  - If there is no free frame, use a page replacement algorithm to select a victim frame
- 3. Write the page in the victim frame to the disk (i.e., page out) if the page is *dirty*
- 4. Read the desired page into the (newly) free frame (i.e., page in)
- 5. Update the page tables
- 6. Restart the instruction that caused the page fault

### Page Replacement



- .Use modify (dirty) bit to reduce overhead of page transfers only modified pages are written to disk
- .Page replacement completes separation between logical memory and physical memory large virtual memory can be provided on a smaller physical memory

## Number of Page Faults vs. Number of Frames




Page faults slow down the system

A good page replacement algorithm should not cause high page faults...

#### FIFO Page Replacement


- Reference string: 7,0,1,2,0,3,0,4,2,3,0,3,0,3,2,1,2,0,1,7,0,1
- Number of frames: 3



- 15 page faults in this case
- FIFO is not always good
  - e.g. first-in != seldom-used
- Suffers from the belady's anomaly

### Belady's Anomaly

#### Page reference string: 1 2 3 4 1 2 5 1 2 3 4 5




Belady's anomaly: page fault rate may increase as the number of allocated frames increases

#### Stack Algorithms

- Can be shown that the set of pages in memory for **n** frames is always a subset of the set of pages that would be in memory with **n+1** frames
- Never exhibit belady's anomaly!
- FIFO is not a stack algorithm, prove it by yourself...
  - You can test the cases of 3 & 4 frames in the previous slide

### Optimal Page Replacement


• Replace the page that will not be used for the longest period of time



- 9 faults in this case
- Has the lowest page fault rate
- Never suffers from the belady's anomaly
- Optimal, but NOT Feasible!

#### LRU (Least Recently Used)

- Replace the page that *has not been used* for the longest period of time
  - Use the recent past as the approximation of the near future




- 12 faults in this case
- Never suffers from the belady's anomaly
- How to implement it?
  - -Clock counters (Timers)
  - -Stack

# LRU Implementation Based on Timer

- Associate with each PTE a time-of-use field (i.e. clock counter)
  - Access a page → update the field with current time
- When page replacement is needed, find the page with the smallest value of the counter

- Problems
  - Requires a search of the PT to find the LRU page
  - Clock counters can overflow

# LRU Implementation Based on Stack



- •Record the access order, instead of the absolute access time
- •Use doubly-linked list to implement the stack
  - because removal from the stack is needed

#### LRU

- Counter and stack implementations are not efficient without special HW support
  - Updating the clock field or the stack must be done for every memory reference
  - Without HW support, use an interrupt for every reference to allow the SW to update the timer/stack → Huge overhead !!!
- Slow even with HW support
  - e.g., Update multiple pointers for each memory reference
- HW support is required, but it should be efficient and easy to implement
  - Several approximation approaches have been proposed

#### Reference bit algorithm

- With each page associate a bit, initially = 0
- When page is referenced, the bit is set to 1
- Replace the page with reference bit = 0 (if one exists)
- However, we do not know the access order
  - No reset!!!! Pages that have not been accessed for a long time cannot be identified...

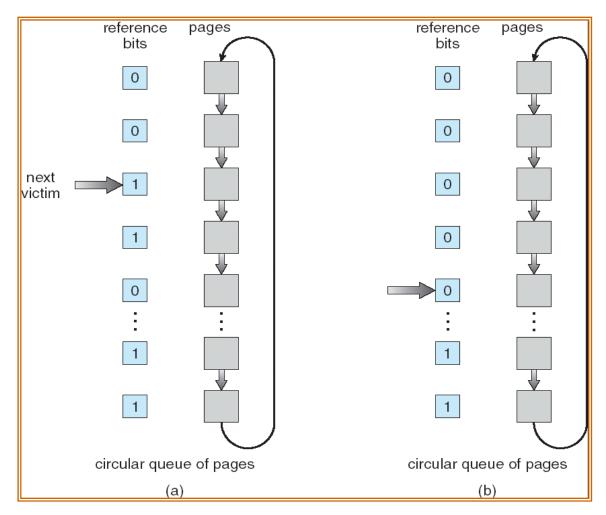
#### Additional reference bit algorithm

- Keep several history bits (e.g., 8 bits) for each page
- Periodically shift the reference bit into the MSB of the history bits
  - Reference bit becomes 0 after each shift
  - History bits of pages that have not been accessed for a long time will become 0...
- Replace the page with the lowest number

- Second chance algorithm (clock algorithm)
  - Need reference bit
  - Clock replacement
    - Scan the PTEs in a clock order
  - If the page has reference bit =  $0 \rightarrow$  replace it
  - If the page has reference bit = 1
    - set reference bit as 0
    - leave the page in memory
    - check next page in clock order
  - If the page is accessed often enough, it will never been replaced

# Second-Chance (Clock) Page-Replacement Algorithm

implemented by using a circular queue



#### Enhanced second chance algorithm

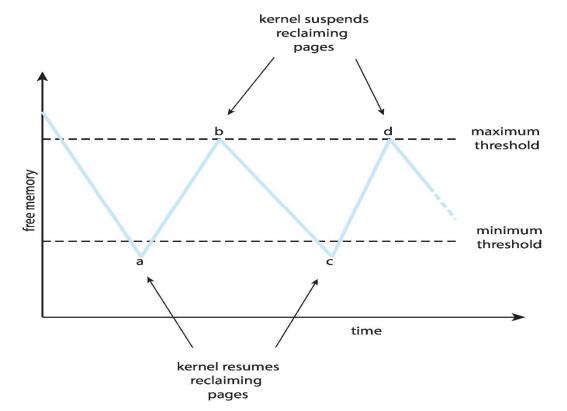
- Consider the pair (reference bit, modify/dirty bit)
   and divide the pages into four classes
  - (0, 0): neither recently used nor modified
    - Best page to replace
    - Why do we check the modify (dirty) bit? **Reduce IO**
  - (0, 1): not recently used but modified
  - (1, 0): recently used but clean
  - (1, 1): recently used and modified
    - Should keep in memory
- Replace the first page encountered in the lowest nonempty class

## **Counting Algorithms**

- Keep a counter of the number of references that have been made to each page
- LFU Algorithm: replaces page with smallest count
  - Since an actively used page should have a large reference count
  - A page used heavily before but never used recently still remains in the memory
    - Aging (i.e., reduce the counts periodically) can be used here
- MFU Algorithm: based on the argument that the page with the smallest count was probably just brought in and has yet to be used

# Page Buffering Algorithms

- Several buffering mechanisms to reduce the page fault time
  - Keep a pool of free frames
    - Fewer page replacement operations in page faults



## Page Buffering Algorithms

- Several buffering mechanisms to reduce the page fault time
  - Other techniques...
    - Write back the modified pages to the paging device when the device is idle (before the pages are selected as victims)
      - Fewer page out operations in page replacement
    - Remember which page was in each frame in the free-frame pool
      - The old page can be reused from the free-frame pool before that frame is rewritten

#### **Allocation of Frames**

- Each process needs a *minimum* number of pages
  - Must at least hold the pages that any single instruction can reference
- Example 1: in a machine that all memory reference instruction have only one memory address ( and no indirection )
  - e.g., LD r1, [100]
  - 2 frames are needed (1 for instruction, 1 for data memory reference)
- Example 2: IBM 370 6 pages to handle SS MOVE instruction:
  - Allow the movement of up to 256 bytes from source to destination
  - instruction is 6 bytes, might span 2 pages
  - 2 pages to handle source
  - 2 pages to handle *destination*
- Example 3: multiple level of indirection
  - LD r1, [[100]] → 3 pages are needed

#### **Allocation of Frames**

- Each process should have a minimum number of pages
- Beyond this number, how many pages a process should have?
  - Equal allocation
  - Proportional allocation
    - Allocate available memory to each process according to its size

$$s_i = \text{size of process } p_i$$
  $s_1 = 10$   
 $S = \sum s_i$   $s_2 = 127$   
 $m = \text{total number of frames}$   $a_1 = \frac{10}{137} \times 64 \approx 5$   
 $a_i = \text{allocation for } p_i = \frac{s_i}{S} \times m$   $a_2 = \frac{127}{137} \times 64 \approx 59$ 

#### Global vs. Local Allocation

- Global replacement algorithms
  - Select a replacement frame from all frames
  - A process cannot control its own page fault rate
- Local replacement algorithms
  - A process select a replacement frame from its own frames
  - Number of frames allocated to a process doesn't change
  - Seldom used frames of the other processes cannot be used by the process
- Global replacement is more common since it generally results in better throughput...

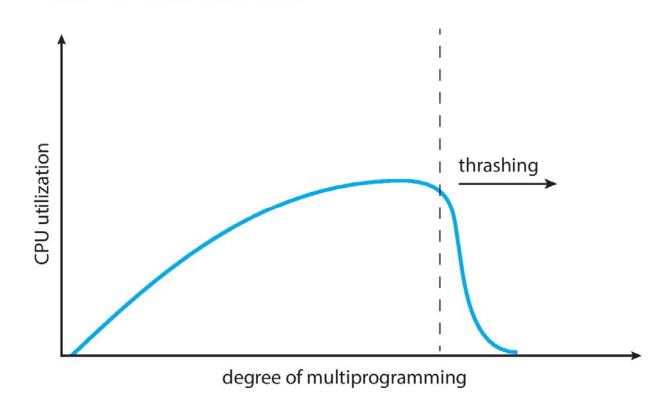
#### **Thrashing**

- If a process does not have "enough" page frames, the page-fault rate will be quite high
  - Since the swapped out pages will be swapped in soon
- In early systems, this led to:
  - low CPU utilization
  - operating system thinks that it needs to increase the degree of multiprogramming
  - another process added to the system, page out some other pages

#### Thrashing

a process is busy swapping pages in and out

# Thrashing (Cont.)



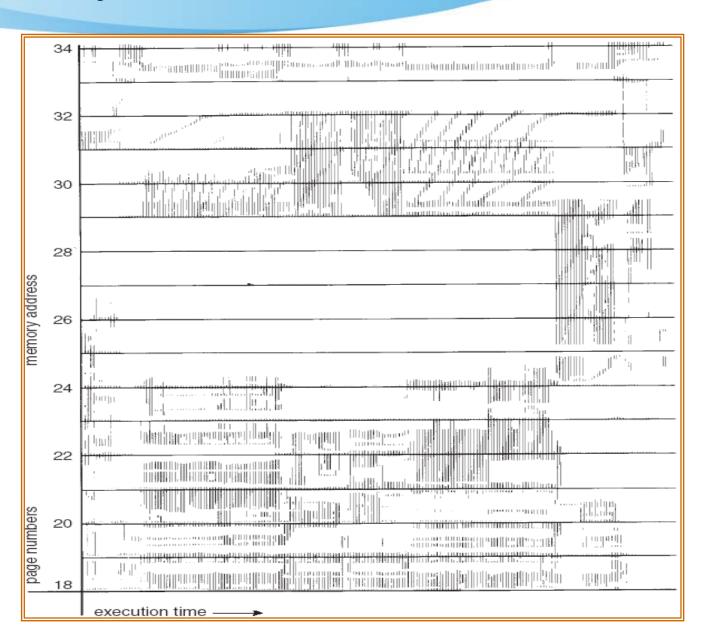
#### Thrashing can be prevented based on

- Working set model, or
- Page fault frequency

#### Locality

- A process should have enough frames to prevent thrashing
- How do we know the number of frames needed by a process?
- The locality model
  - As a process executes, it moves from locality to locality
  - A locality: set of pages that are actively used together
    - E.g., a function call changes from one locality to another
- Allocate frames to accommodate the current locality of each process

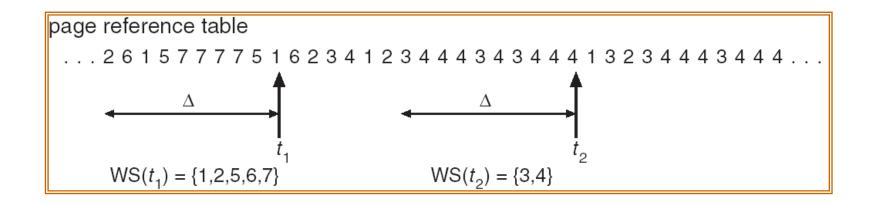
## Memory Reference Pattern



# **Working-Set Model**

- Based on the assumption of locality
- $\Delta \equiv$  working-set window  $\equiv$  a fixed number of page references Example: 10,000 references
- $WSS_i$  (working set size of Process  $P_i$ ) = total number of pages referenced in the most recent  $\Delta$  (varies in time)
  - if  $\Delta$  too small  $\rightarrow$  will not encompass entire locality
  - if  $\Delta$  too large  $\rightarrow$  will encompass several localities
  - if  $\Delta$  = ∞ → will encompass entire program
- Allocate at least WSSi pages for Pi
- $D = \sum WSS_i \equiv \text{total demand frames}$
- if D > m (total # of available frames)  $\Rightarrow$  Thrashing
  - if D > m, suspend processes 1-by-1 until  $D \le m$

# Working-Set Model



Tracking the WSS at each memory reference leads to a huge overhead

Implement the WS model by using timer interrupt & reference bit

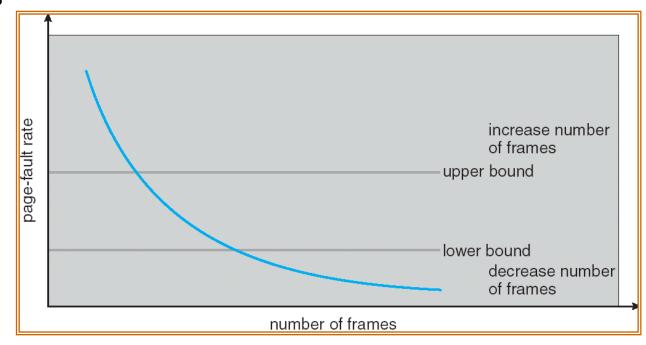
- just an approximation

# Page-Fault Frequency Scheme

- Another way to prevent thrashing
- Establish "acceptable" page-fault rate
  - If actual rate too low, process loses frame
  - If actual rate too high, process gains frame

If not all the process have acceptable rates → suspend a

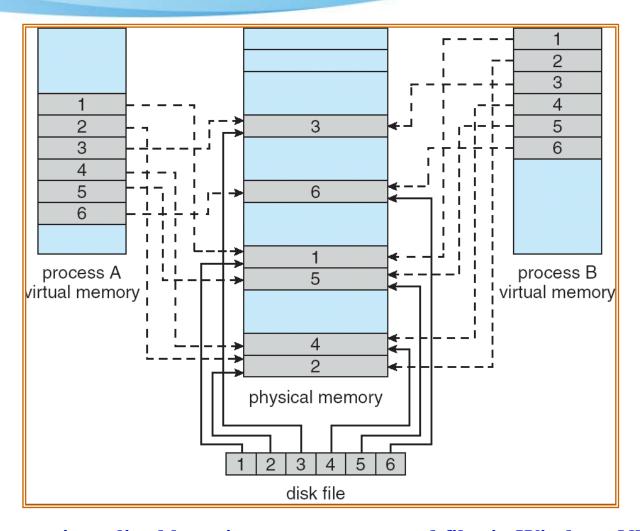
process



#### **Memory-Mapped Files**

- Memory-mapped file I/O allows file I/O to be treated as memory access by mapping a disk block to a page in memory
  - Usually, using the mmap() system call
- A file is read using *demand paging*. A page-sized portion of the file is read from the file system into a physical page. Subsequent reads/writes from/to the file are treated as ordinary memory accesses.
- Avoid read()/write() system calls during file access
  - Reduce overhead
- Also allows several processes to map the same file → allowing the pages in memory to be shared

# **Memory Mapped Files**

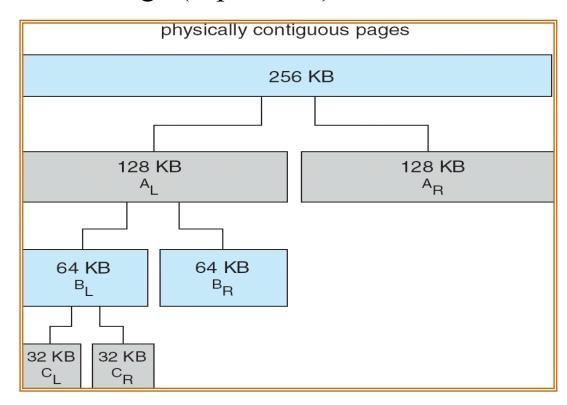


Shared memory is realized by using memory mapped files in Windows NT,2000, XP

# **Allocating Kernel Memory**

#### Buddy System

- Power-of-2 allocator
- Split (if necessary) when allocation
- Merge (if possible) when free



- Advantage
  - ✓ quick alloc/free
- Disadvantage
  - ✓ fragmentation

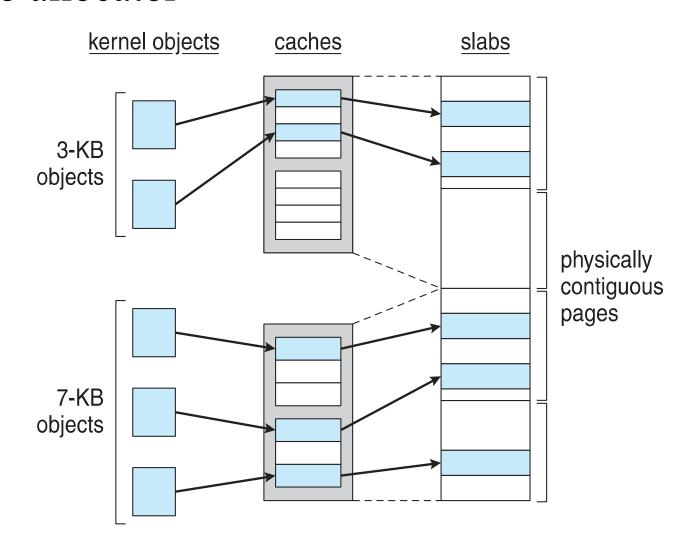
## **Allocating Kernel Memory**

#### Slab Allocator

- Slab: one or more physically contiguous pages
- Cache: one or more slabs
- a cache for each unique kernel data structure
  - Each cache filled with a type of **objects** (data structures)
- When cache created, filled with free objects
- When structures stored, objects marked as used
- If slab is full of used objects, use a new empty slab for the coming requests
- Benefits: no fragmentation, quick allocation/free

# **Allocating Kernel Memory**

#### • Slab allocator



#### Slab Allocator in Linux

- PCB in Linux: struct task struct
- Approx 1.7KB of memory
- New task -> allocate new struct from cache
- Slab can be in three possible states
  - 1. Full all used
  - 2. Empty all free
  - 3. Partial mix of free and used
- Upon request, slab allocator
  - 1. Uses free struct in partial slab
  - 2. If none, takes one from empty slab
  - 3. If no empty slab, create new empty

# Other Considerations of a Paging System

- Major decisions of a paging system
  - Page replacement algorithm
  - Page allocation policy
- Other considerations
  - Prepaging
  - Page size
  - TLB reach
  - Inverted Page Tables
  - Program structure
  - IO interlock

#### **Prepaging**

- To reduce the large number of page faults occur at process startup
- Prepage some of the pages a process will need,
   before they are referenced
- But if prepaged pages are unused, I/O and memory was wasted
- Assume s pages are prepaged and a fraction α of the pages is used
  - Is cost of  $s * \alpha$  saved pages faults > or < the cost of prepaging  $s * (1-\alpha)$  unnecessary pages?
  - $-\alpha$  near zero  $\Rightarrow$  prepaging loses

#### **TLB Reach**

- TLB Reach The amount of memory accessible from the TLB
- TLB Reach = (Number of TLB Entries) x (Page Size)
- Ideally, the working set of each process is stored in the TLB. Otherwise, a lot of TLB misses and page table accesses.
- Increase the page size can increase TLB reach
- However, this may lead to an increase in fragmentation as not all applications require a large page size
- Provide multiple page sizes
  - allows applications that require larger page sizes the opportunity to use them without an increase in fragmentation
  - Solaris on UltraSPARC uses both 8KB and 4MB page sizes

#### Page Size

- Page size selection must take into consideration
  - Large page size, good for
    - page table size
    - page fault #
    - I/O overhead
      - Transferring a large page is more cost effective
    - TLB reach
  - Small page size, good for
    - locality
    - fragmentation

#### **Program Structure**

#### Program structure

- Int[128,128] data;
- Each row is stored in one page
- OS allocates fewer than 128 frames to this process
- Program 1

 $128 \times 128 = 16,384$  page faults

- Program 2

for 
$$(i = 0; i < 128; i++)$$
  
for  $(j = 0; j < 128; j++)$   
data $[i,j] = 0;$ 

128 page faults

#### **Data structures**

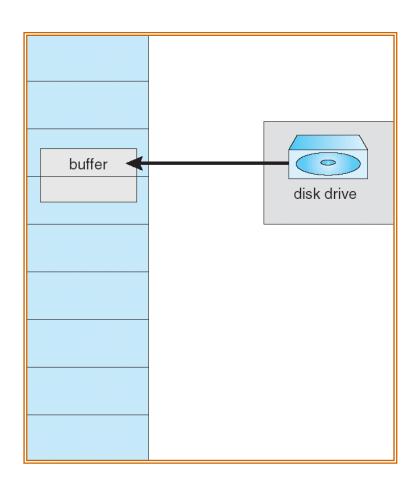
-pointers, hash → poor locality -stack → good locality

# **Inverted Page Table (IPT)**

- Reduce the memory requirement of the page table
- However, IPT does not contain the complete information about the VAS of a process
  - No information about the not-in-memory pages!!
    - E.g., where the page is (in the swap area), the protection bits...
- Therefore, an external page table (one per process) must be kept
  - The format just like the traditional page tables
  - These tables are referenced only when a page fault occurs
  - These tables are themselves paged in/out
  - A page fault may cause the VM manager to load a page of the external page table into memory (an extra I/O operation)

#### I/O Interlock

- I/O Interlock pages must sometimes be *locked* in the memory
  - e.g. pages that are used for copying a file from a device must be locked from being selected for eviction by a page replacement algorithm.
- Each page is associated with a *lock* bit



# Operating System Examples -

Windows XP

• Solaris

#### Windows XP

- Uses demand paging with **clustering**. Clustering brings in pages surrounding the faulting page.
- Processes are assigned working set minimum and working set maximum
- Working set minimum is the minimum number of pages the process is guaranteed to have in memory
- A process may be assigned as many pages up to its working set maximum
- When the amount of free memory in the system falls below a threshold, **automatic working set trimming** is performed to restore the amount of free memory
- Working set trimming removes pages from processes that have pages in excess of their working set minimum

#### **Solaris**

- Maintains a list of free pages to assign to faulting processes
- Lotsfree threshold parameter (amount of free memory) to begin paging
- *Desfree* threshold parameter to increasing paging
- *Minfree* threshold parameter to being swapping
- Paging is performed by *pageout* process
- Pageout process scans pages using a modified clock algorithm
- *Scanrate* is the rate at which pages are scanned. This ranges from *slowscan* to *fastscan*
- Pageout is called more frequently depending upon the amount of free memory available

# Solaris 2 Page Scanner

